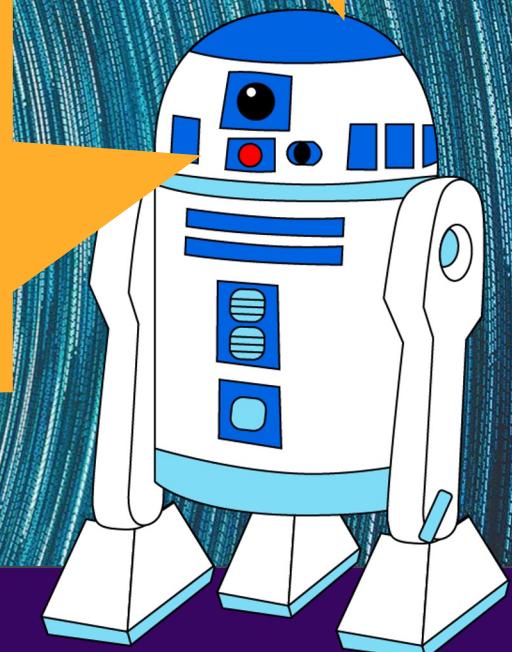


CIS 4210/5210:
ARTIFICIAL INTELLIGENCE

Informed Search

HW2 is due Wednesday by 11:59pm.

If you need an extension, you must request it 24 hours in advance to avoid a late penalty.



Review: Search problem definition

States: a set S

An *initial state* $s_i \in S$

Actions: a set A

$\forall s \text{ } Actions(s) = \text{the set of actions that can be executed in } s, \text{ that are applicable in } s.$

Transition Model: $\forall s \forall a \in Actions(s) \text{ } Result(s, a) \rightarrow s_r$

s_r is called a *successor* of s

$\{s_i\} \cup Successors(s_i)^*$ = *state space*

Path cost (*Performance Measure*): Must be additive

e.g. sum of distances, number of actions executed, ...

$c(x, a, y)$ is the step cost, assumed ≥ 0

(where action a goes from state x to state y)

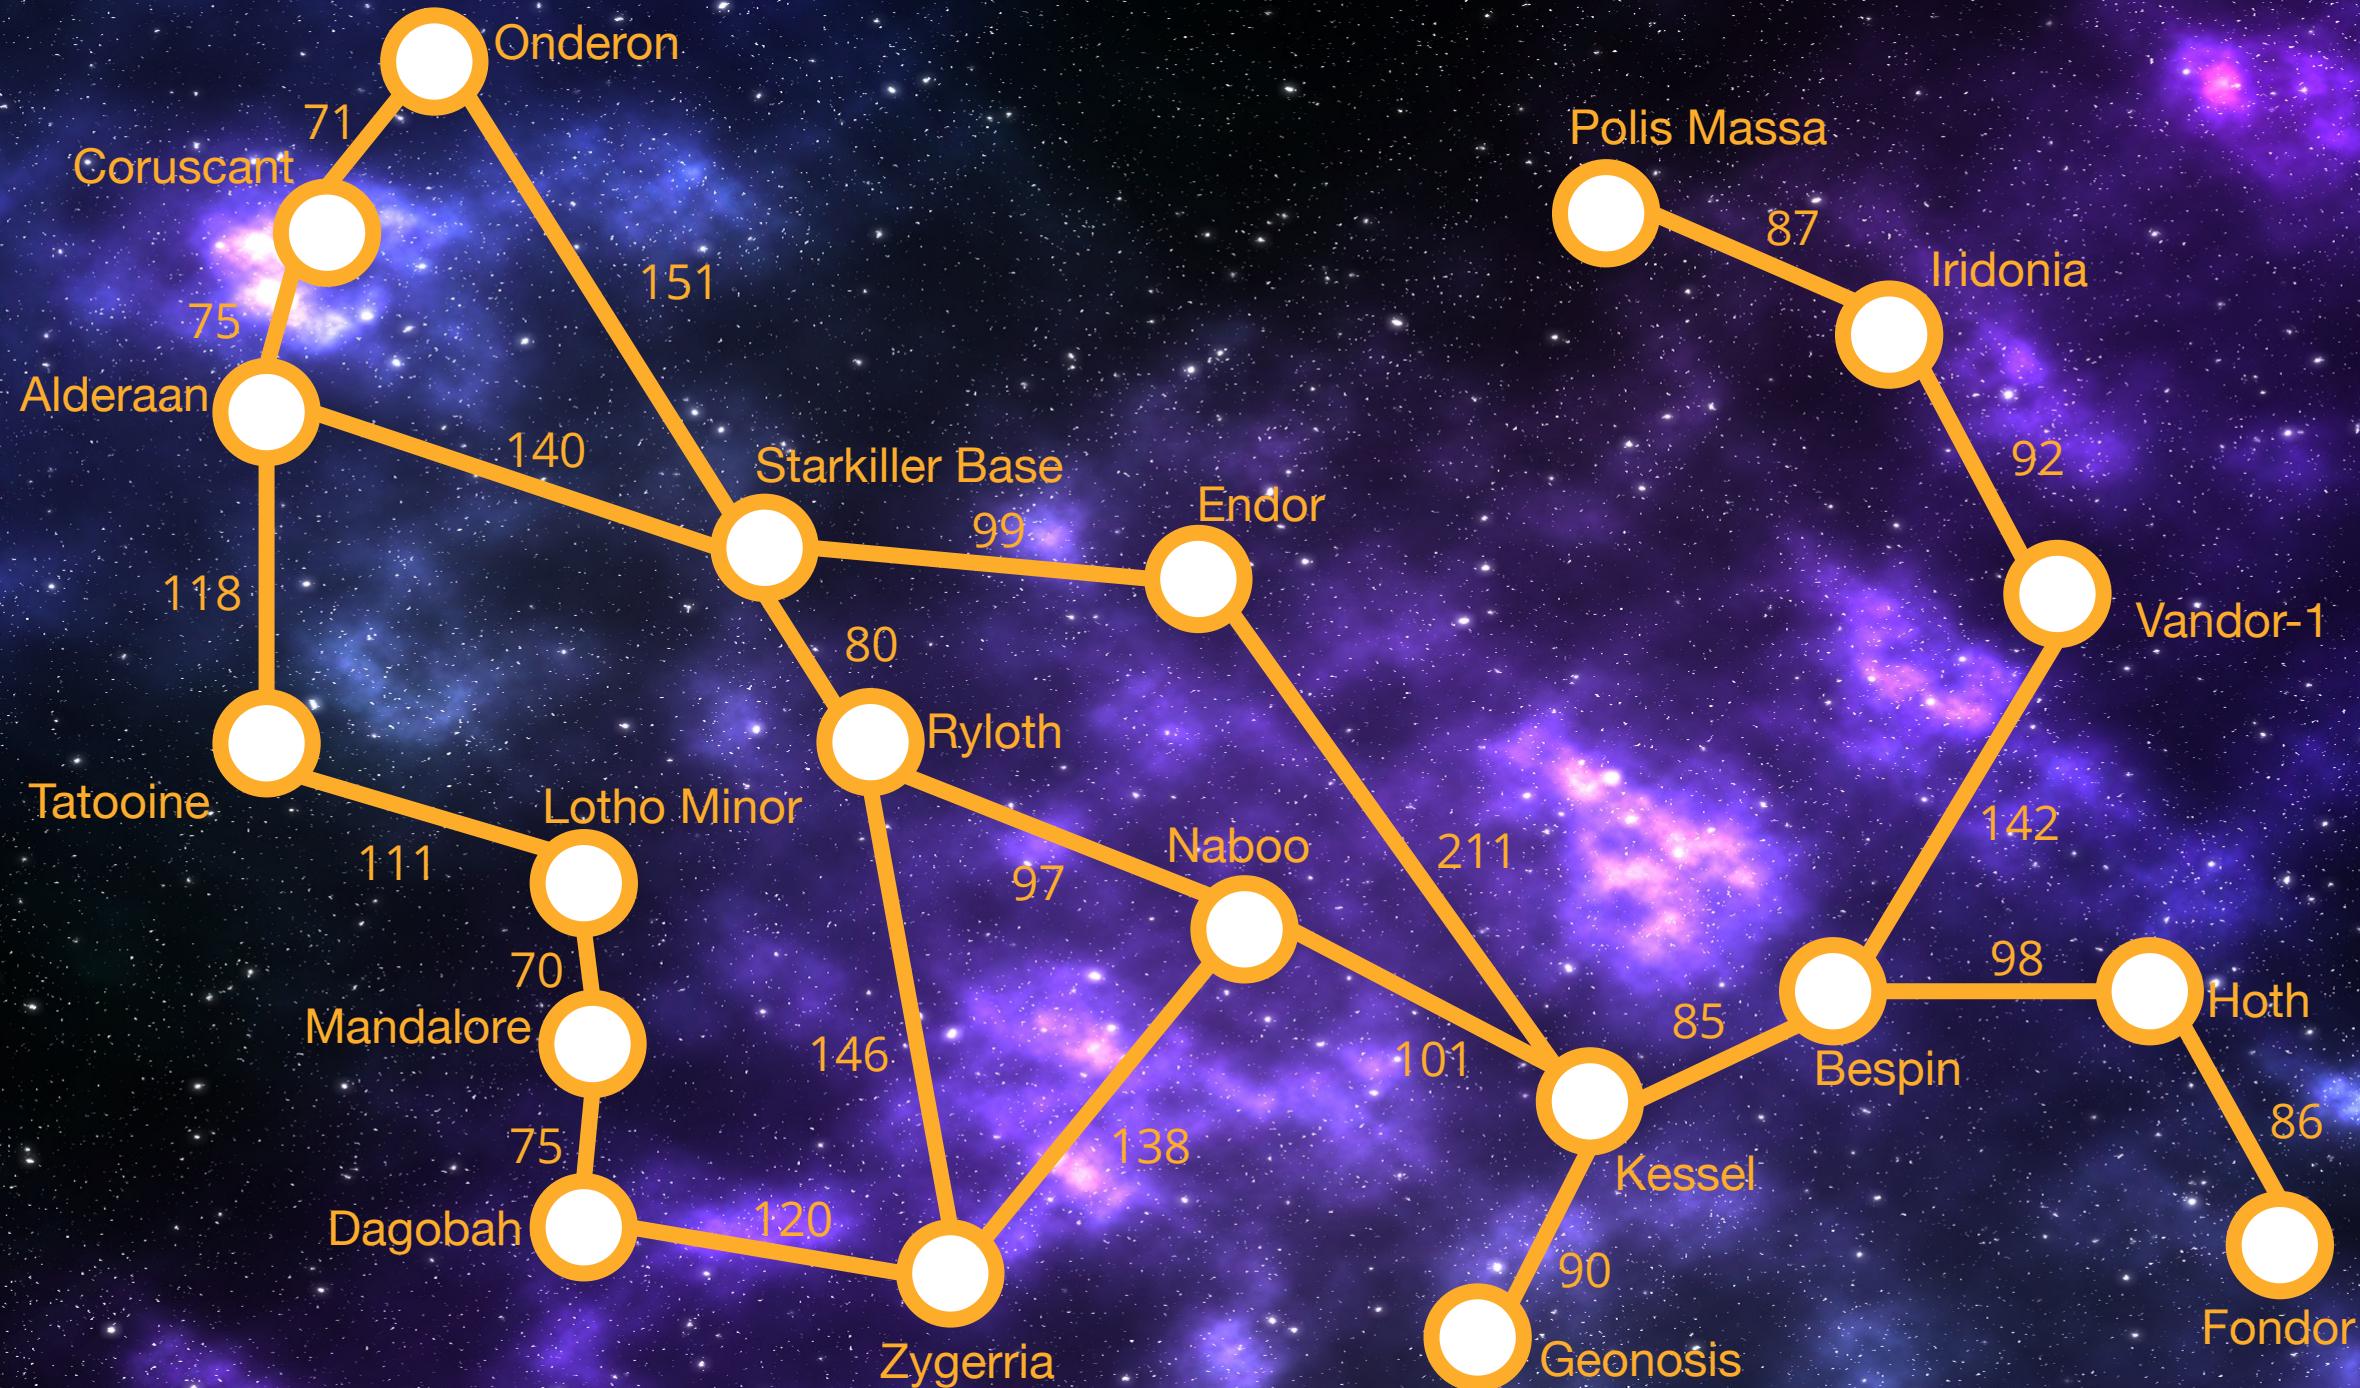
Goal test: $Goal(s)$

Can be implicit, e.g. $checkmate(s)$

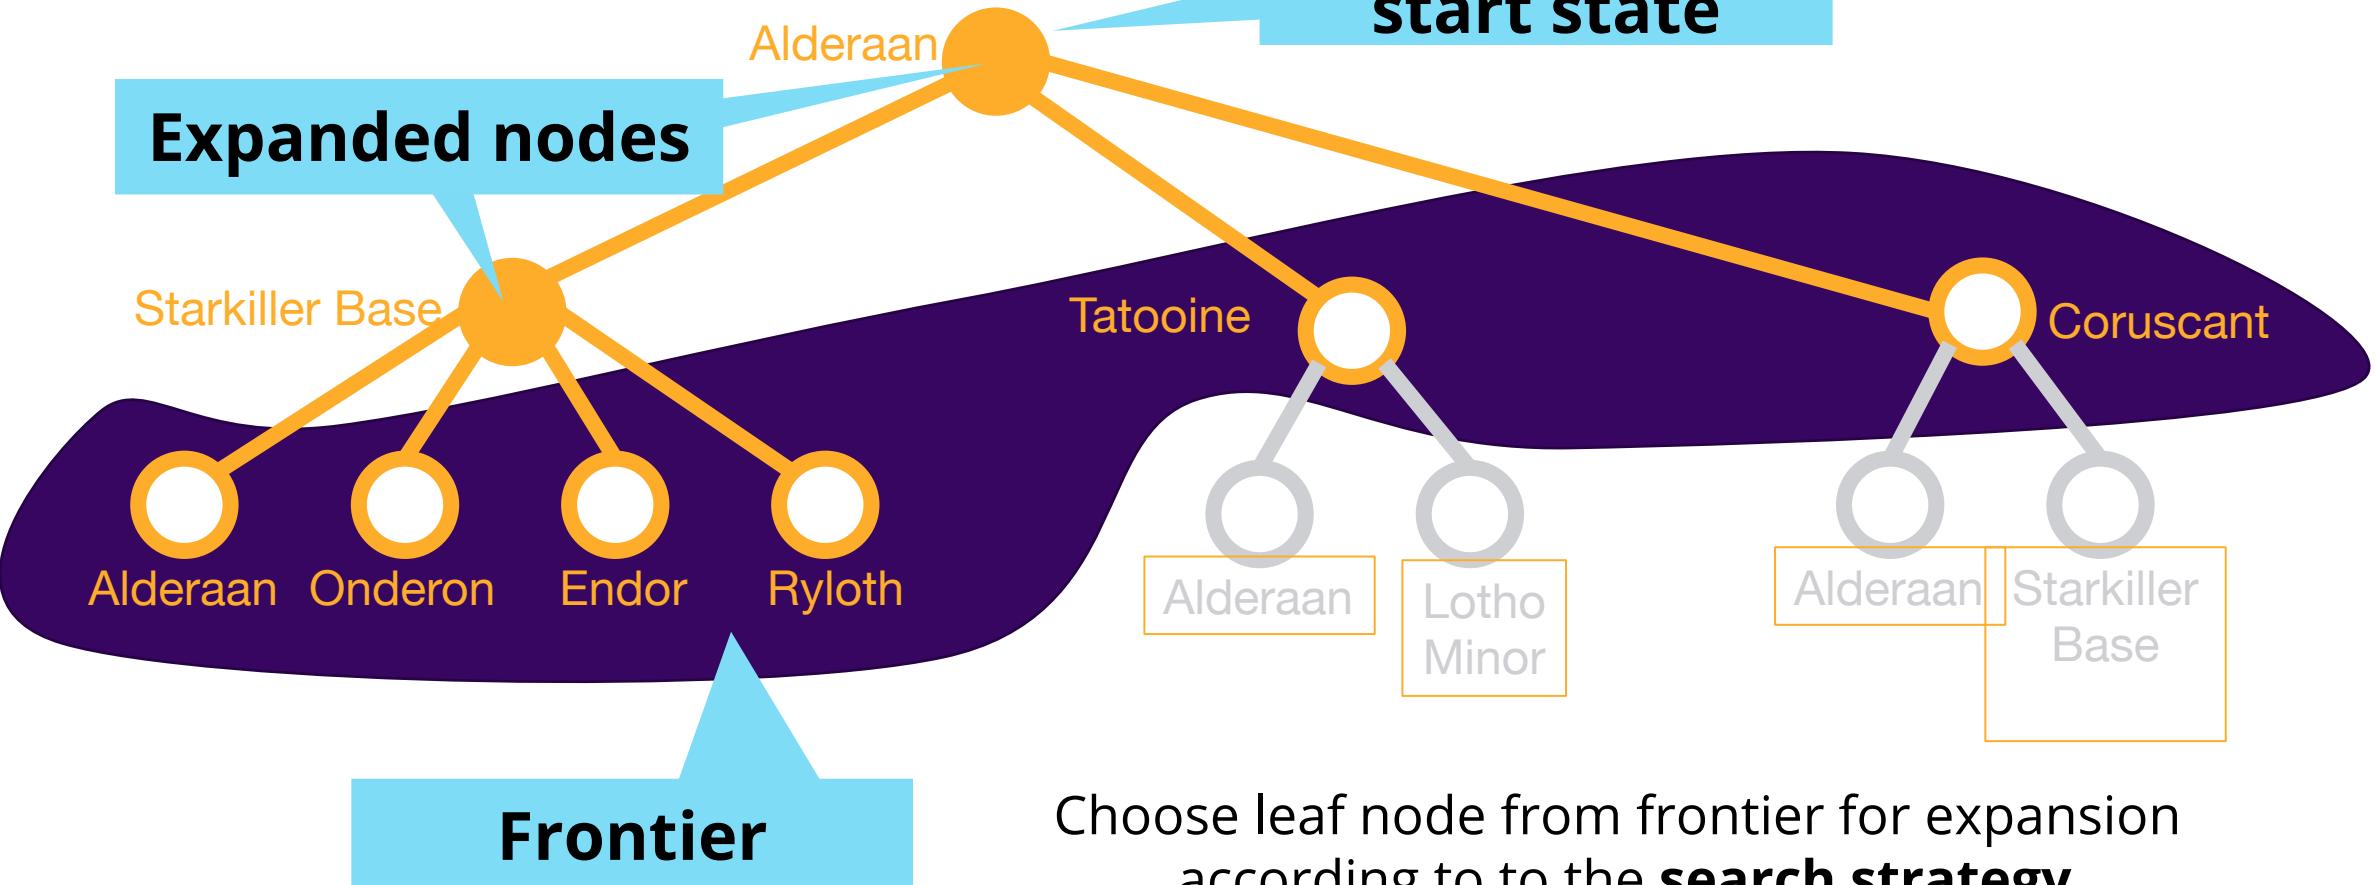
s is a *goal state* if $Goal(s)$ is true

Review: Useful Concepts

- *State space*: the set of all states reachable from the initial state by *any* sequence of actions
 - *When several operators can apply to each state, this gets large very quickly*
 - *Might be a proper subset of the set of configurations*
- *Path*: a sequence of actions leading from one state s_j to another state s_k
- *Frontier*: those states that are available for *expanding* (for applying legal actions to)
- *Solution*: a path from the initial state s_i to a state s_g that satisfies the goal test



Search Tree



Review: Search Strategies

Strategy = order of tree expansion

- Implemented by different **queue structures** (LIFO, FIFO, priority)

Dimensions for evaluation

- *Completeness*- always find the solution?
- *Optimality* - finds a least cost solution (lowest path cost) **first?**
- *Time complexity* - # of nodes generated (*worst case*)
- *Space complexity* - # of nodes simultaneously in memory (*worst case*)

Time/space complexity variables

- b , *maximum branching factor* of search tree
- d , *depth* of the shallowest goal node
- m , maximum length of any path in the state space (potentially ∞)

Animation of Graph BFS algorithm
set to music 'flight of bumble bee'

<https://youtu.be/x-VTfcmrLEQ>

Animation of Graph DFS algorithm

Depth First Search of Graph

set to music 'flight of bumble bee'

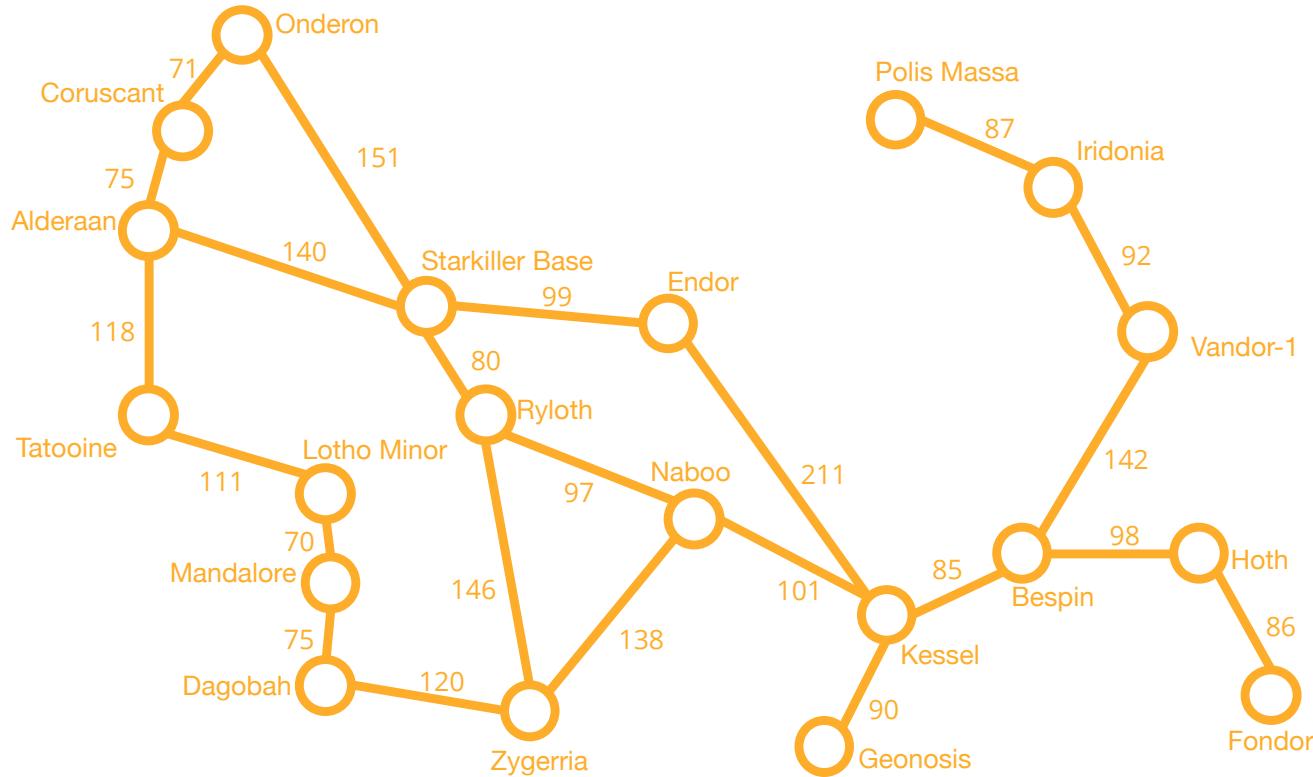
“Uniform Cost” Search

“In computer science, *uniform-cost* search (UCS) is a tree search algorithm used for traversing or searching a *weighted* tree, tree structure, or graph.” - Wikipedia

Motivation: Map Navigation Problems

All our search methods so far
assume *step-cost* = 1

This is only true for some problems



$g(N)$: the path cost function

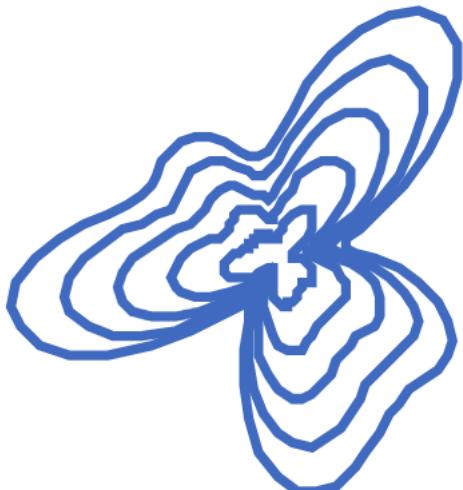
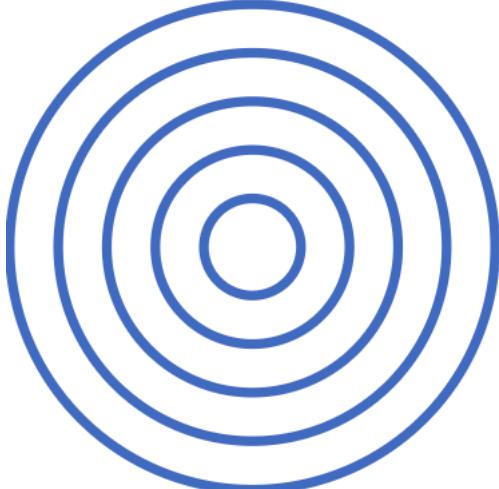
- Our assumption so far: All moves equal in cost
 - Cost = # of nodes in path-1
 - $g(N) = \text{depth}(N)$ in the search tree
- More general: Assigning a (potentially) unique cost to each step
 - N_0, N_1, N_2, N_3 = nodes visited on path p from N_0 to N_3
 - $C(i,j)$: Cost of going from N_i to N_j
 - If N_0 the root of the search tree,
$$g(N_3) = C(0,1) + C(1,2) + C(2,3)$$

Uniform-cost search (UCS)

- Extension of BF-search:
 - **Expand node with *lowest path cost***
- Implementation:
 - *frontier* = priority queue ordered by $g(n)$
- Subtle but significant difference from BFS:
 - Tests if a node is a goal state when it is selected for expansion, **not when it is added to the frontier.**
 - Updates a node on the frontier if a better path to the same state is found.
 - So always enqueues a node *before checking whether it is a goal.*

WHY???

Shape of Search



- **Breadth First Search** explores equally in all directions. Its frontier is implemented as a FIFO queue. This results in smooth contours or “plies”.
- **Uniform Cost Search** lets us prioritize which paths to explore. Instead of exploring all possible paths equally, it favors lower cost paths. Its frontier is a priority queue. This results in “cost contours”.

This may remind you of something...



Which algorithm from your algorithms class does Uniform Cost Search remind you of?

PollEv.com/ccb

Uniform Cost Search (UCS) vs. Dijkstra's algorithm

Foundational Concepts:

- Dijkstra's: Shortest path from initial to all vertices in a weighted graph.
- UCS: AI's pathfinding. Shortest path from initial state to a goal state.

Similarities:

- Use of **Priority Queue** to pick lowest cost node.
- **Greedy Nature**: Always choose optimal local solution.
- Both operate on a **weighted graph**.
- Guarantee **Optimality** for non-negative weights.

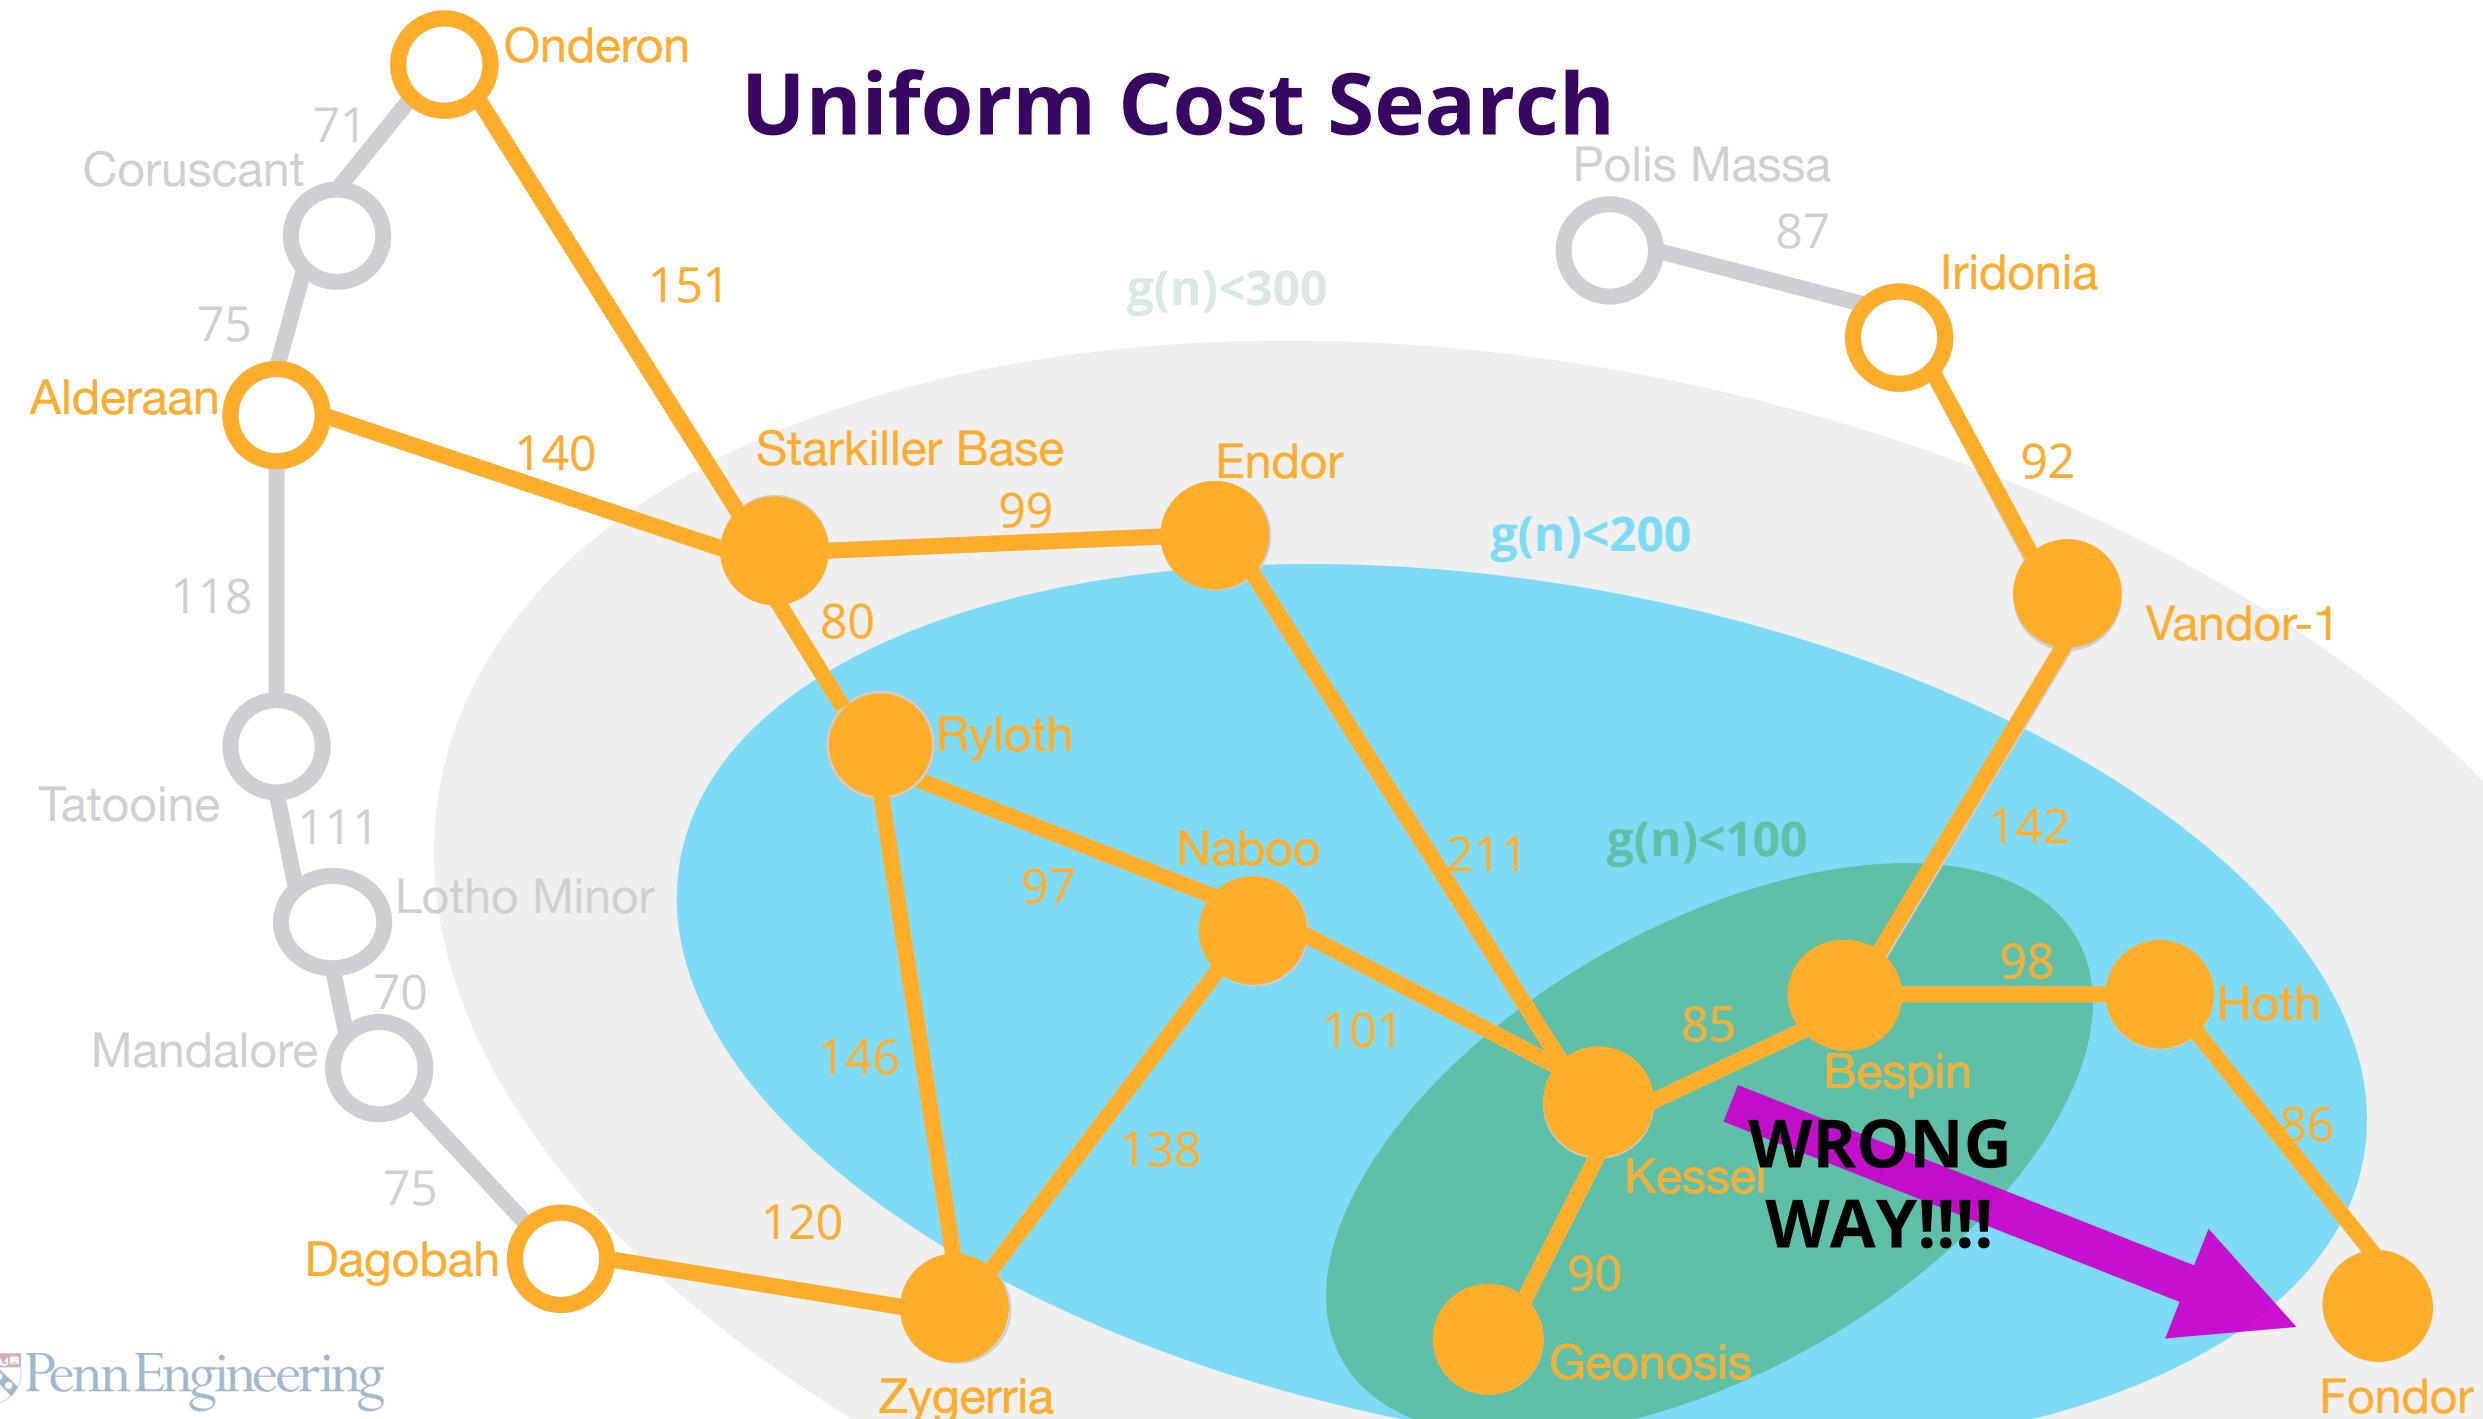
Differences:

- **Objective**: UCS targets specific goal node; Dijkstra's finds paths to all nodes.
- **Termination**: UCS stops at goal; Dijkstra's stops after all nodes have been reached.

Relation:

- UCS is a "targeted" Dijkstra's: it seeks a path to a specific node, not all nodes.

Uniform Cost Search



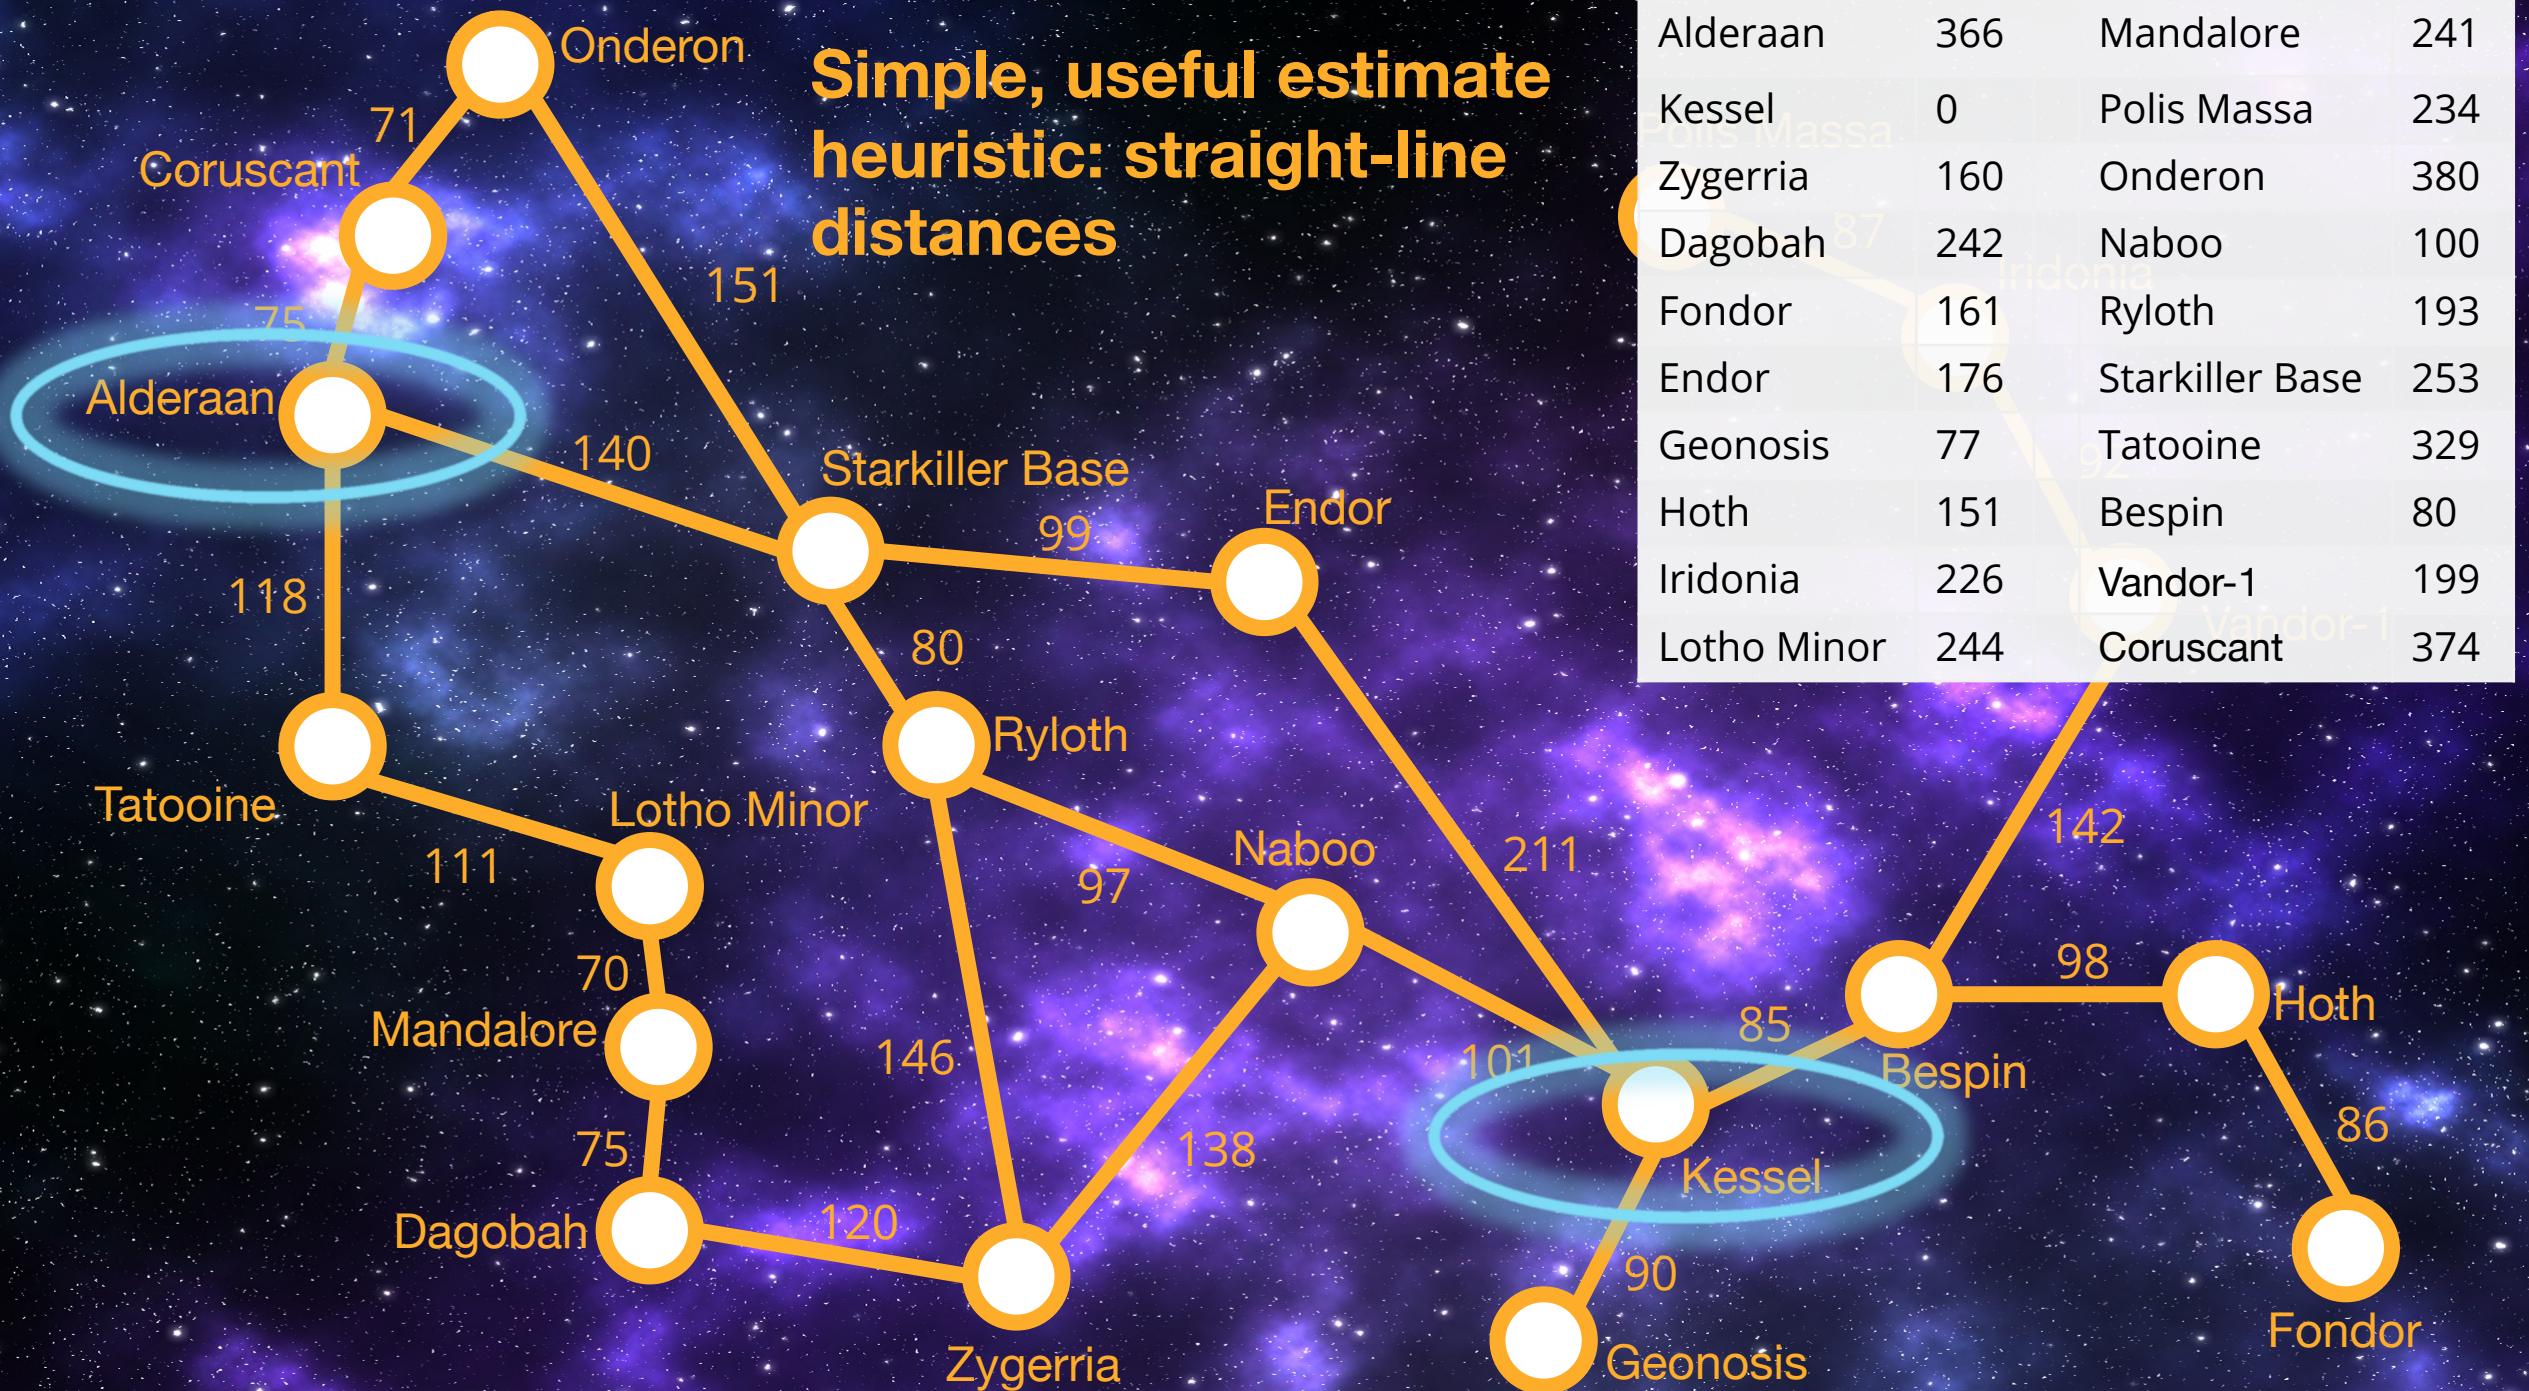
A better idea than UCS...

- Node expansion based on *an estimate* which *includes distance to the goal*
- General approach of **informed search**:
 - **Best-first search**: node selected for expansion based on an *evaluation function $f(n)$*
 - ✓ $f(n)$ includes estimate of distance to goal (*new idea!*)
- Implementation: **Sort our frontier queue** by this **new $f(n)$** .

Instances of best-first search:

- **greedy search**
- **A^* search**

Simple, useful estimate heuristic: straight-line distances



Heuristic (estimate) functions



Heureka! ---Archimedes

[dictionary] “**A rule of thumb, simplification, or educated guess that reduces or limits the search for solutions in domains that are difficult and poorly understood.**”

Heuristic knowledge is useful, but not necessarily correct.

Heuristic algorithms use heuristic knowledge to solve a problem.

A *heuristic function* $h(n)$ takes a state n and returns an *estimate* of the distance from n to the goal.

Greedy Best-First Search

First attempt at integrating heuristic knowledge

Best-first search

Basic idea:

Select node for expansion with minimal *evaluation function $f(n)$*

- where $f(n)$ is some function that includes *estimate heuristic $h(n)$* of the remaining distance to goal

Implement using priority queue

Exactly UCS with $f(n)$ replacing $g(n)$

Greedy best-first search: $f(n) = h(n)$

Expands the node that *is estimated* to be closest to goal

Completely ignores $g(n)$: the cost to get to n

In our Romanian map, $h(n) = h_{SLD}(n)$ = straight-line distance from n to Bucharest

In a grid, the heuristic distance can be calculated using the “Manhattan distance”:

```
def heuristic(a, b):
    # Manhattan distance on a square grid
    return abs(a.x - b.x) + abs(a.y - b.y)
```

Greedy best-first search

```
frontier = PriorityQueue()
frontier.put(start, 0)
came_from = {}
came_from[start] = None

while not frontier.empty():
    current = frontier.get()

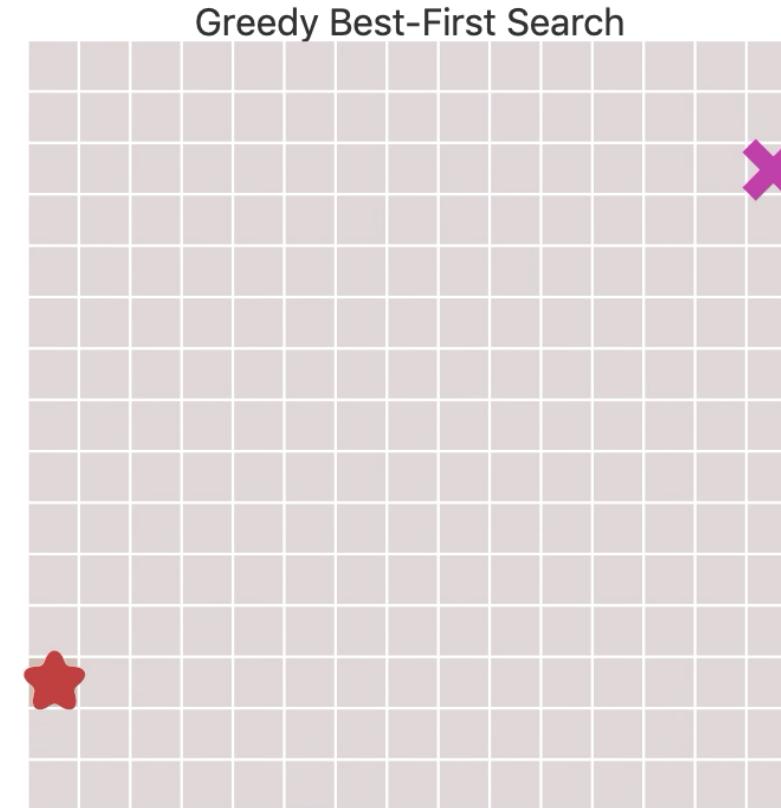
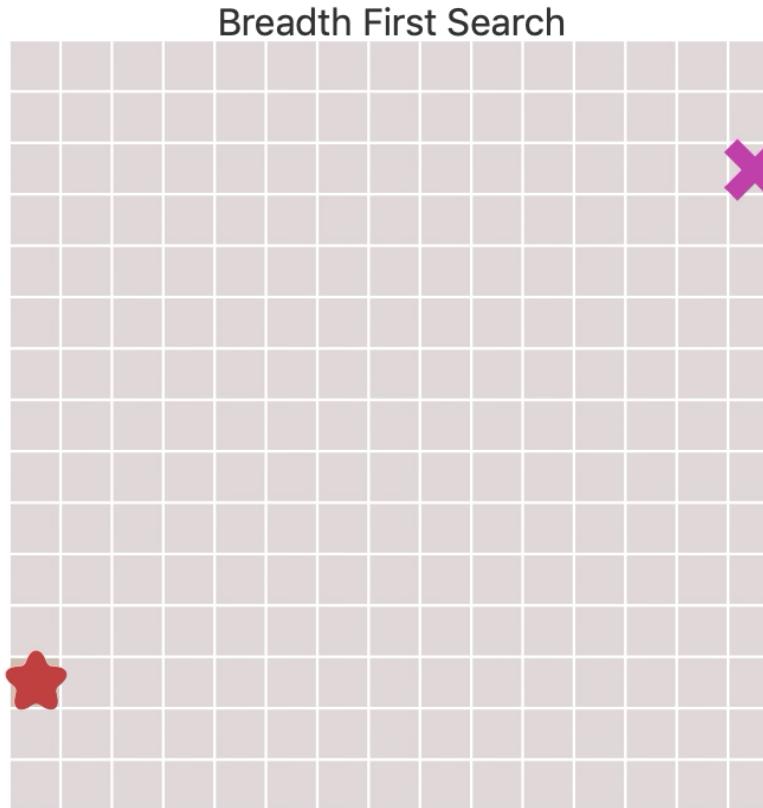
    if current == goal:
        break

    for next in graph.neighbors(current):
        if next not in came_from:
            priority = heuristic(goal, next)
            frontier.put(next, priority)
            came_from[next] = current
```

Code from Amit Patel
of Red Blob Games

This is not the solution
to the HW assignment.

BFS v. Greedy Best-First Search

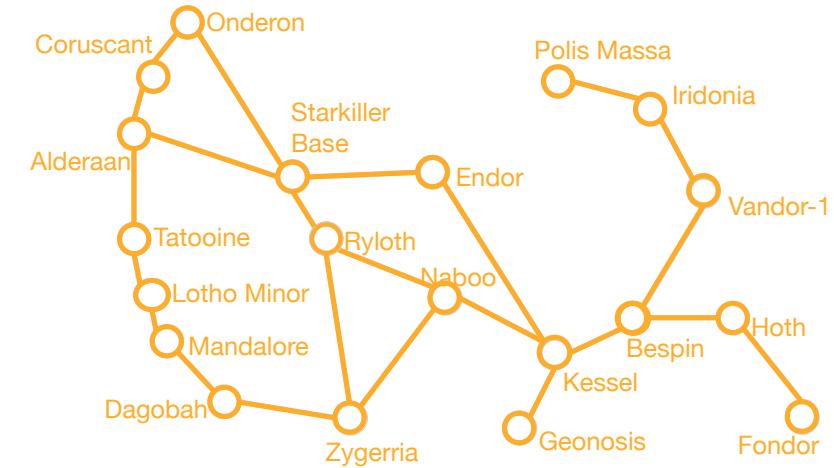


<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Greedy best-first search example

Initial State: **Alderaan**
Goal State: **Kessel**

Alderaan



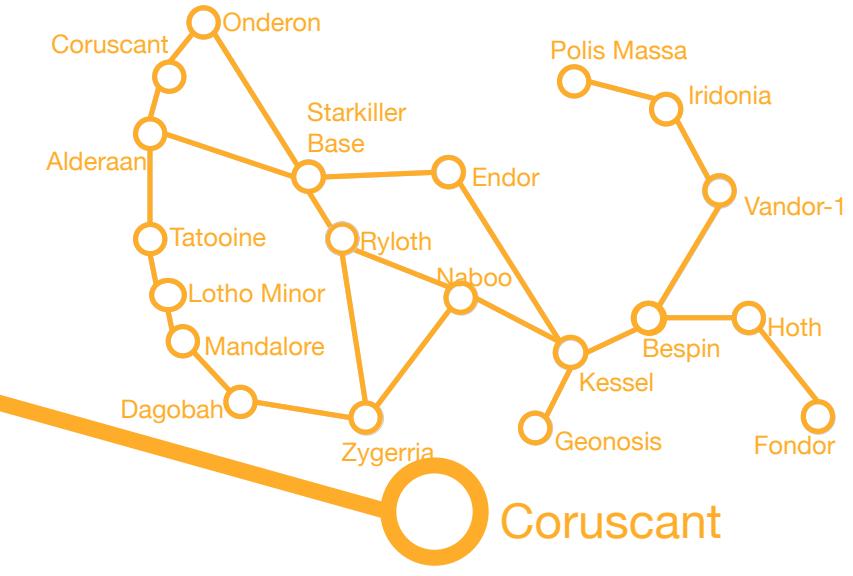
Frontier queue:
Alderon 366

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Greedy best-first search example

Initial State: Alderaan

Goal State: Kessel



Frontier queue:

Starkiller Base 253

Tatooine 329

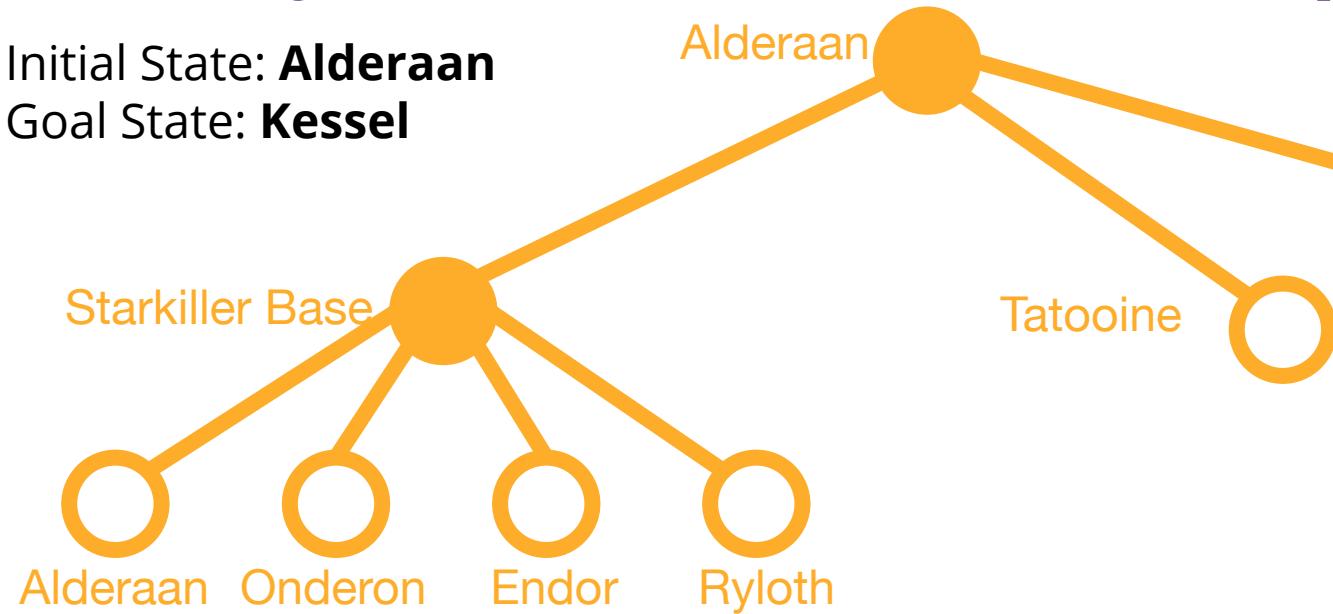
Coruscant 374

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Greedy best-first search example

Initial State: Alderaan

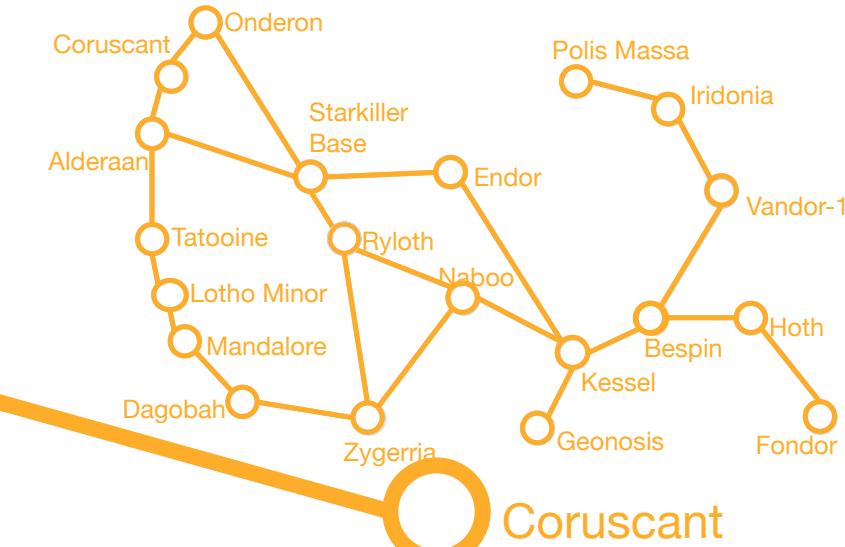
Goal State: Kessel



Are we doing **graph**
search or **tree** search?

Frontier queue:

Tatooine	329
Coruscant	374
Alderaan	366
Onderon	380
Endor	176
RylOTH	193



Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Graph search versus tree search



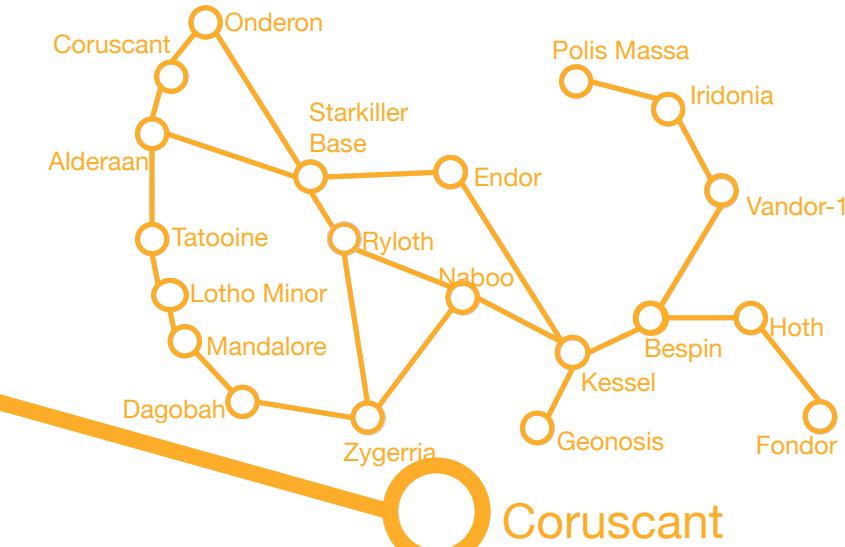
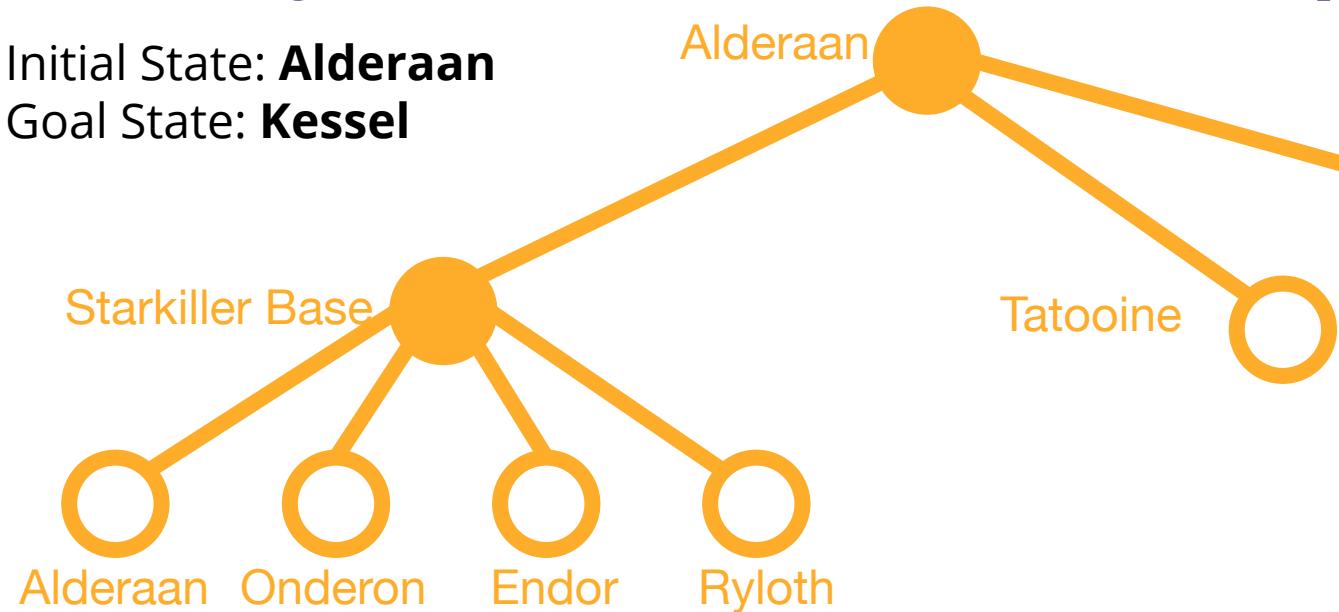
What is the difference between graph search and tree search?

PollEv.com/ccb

Greedy best-first search example

Initial State: Alderaan

Goal State: Kessel



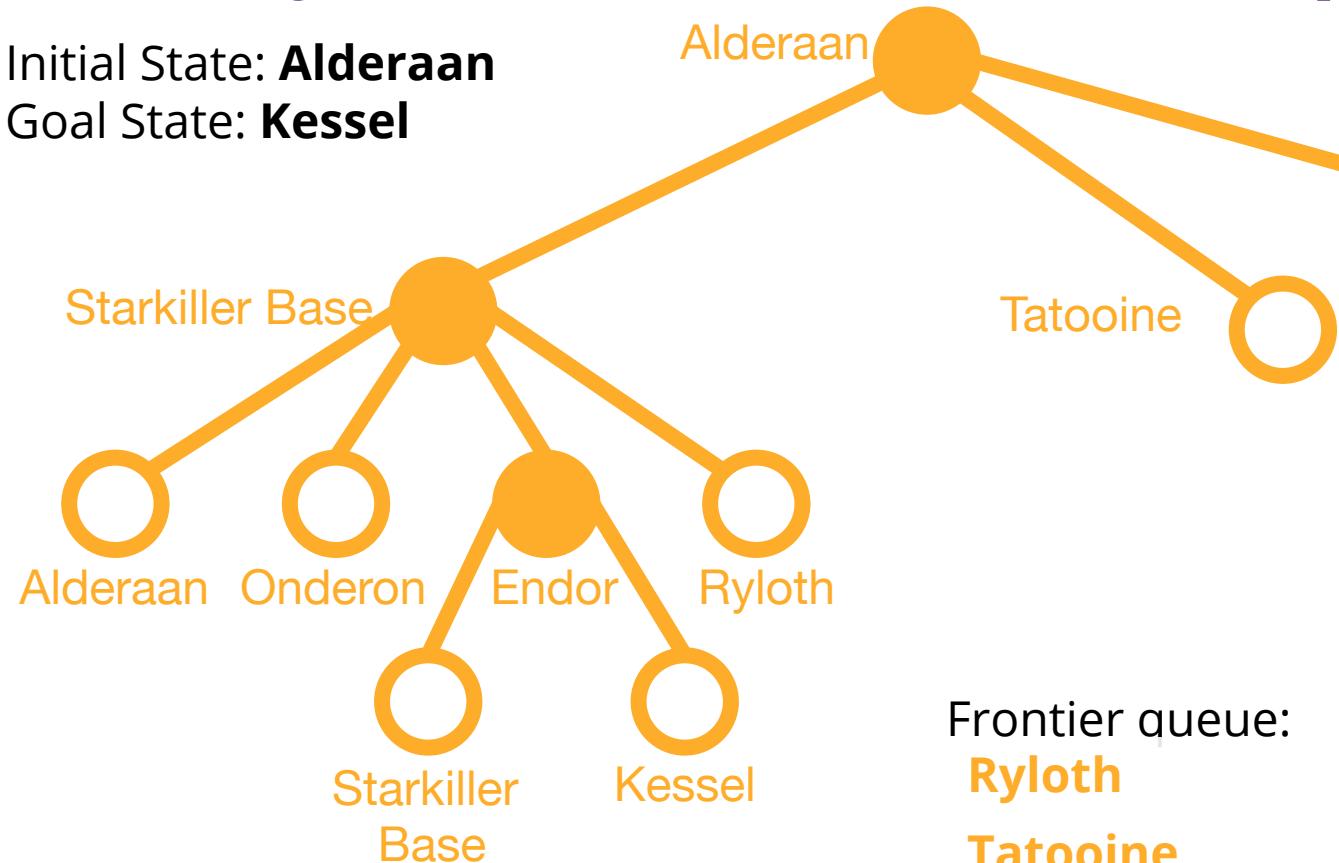
Endor	176
RylOTH	193
Tatooine	329
Alderaan	366
Coruscant	374
Onderon	380

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	RylOTH	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Greedy best-first search example

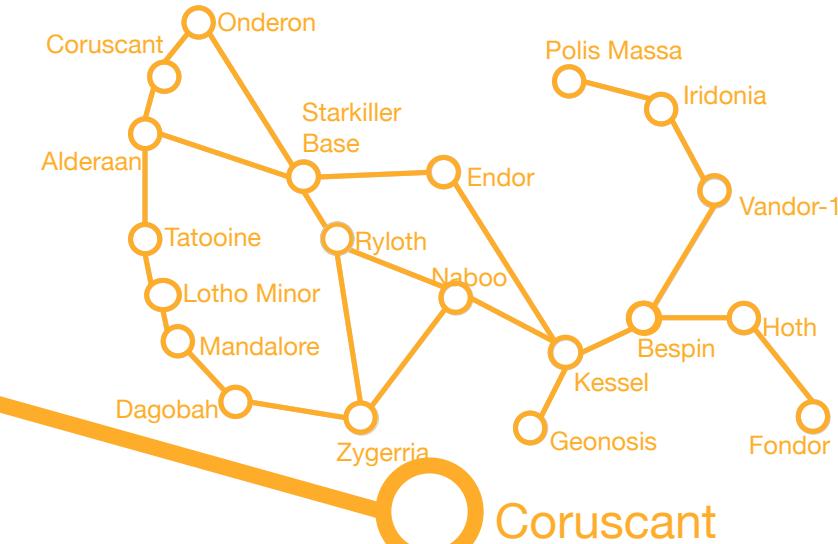
Initial State: Alderaan

Goal State: Kessel



Frontier queue:

RylOTH	193
Tatooine	329
Alderaan	366
Coruscant	374
Onderon	380
Starkiller Base	253
Kessel	0

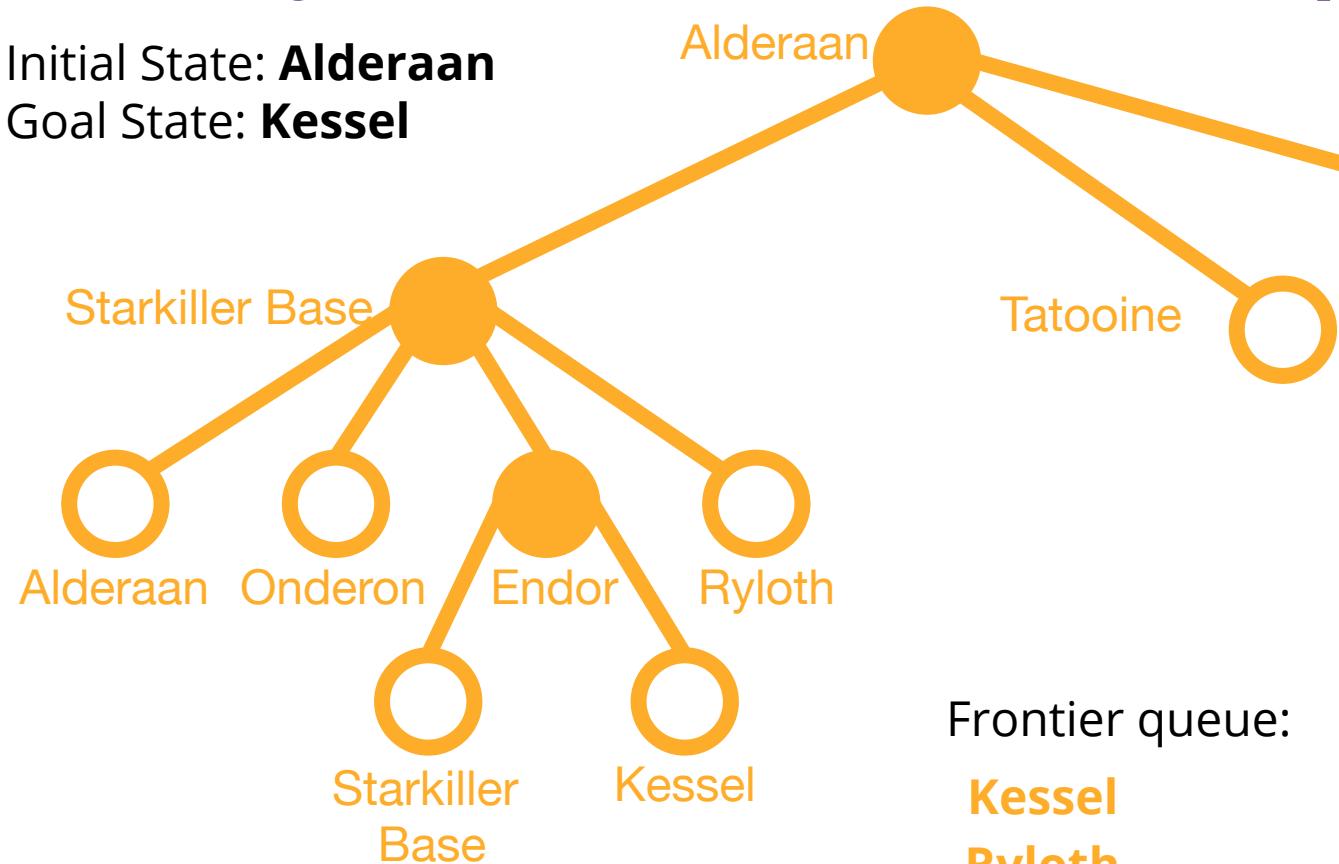


Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	RylOTH	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

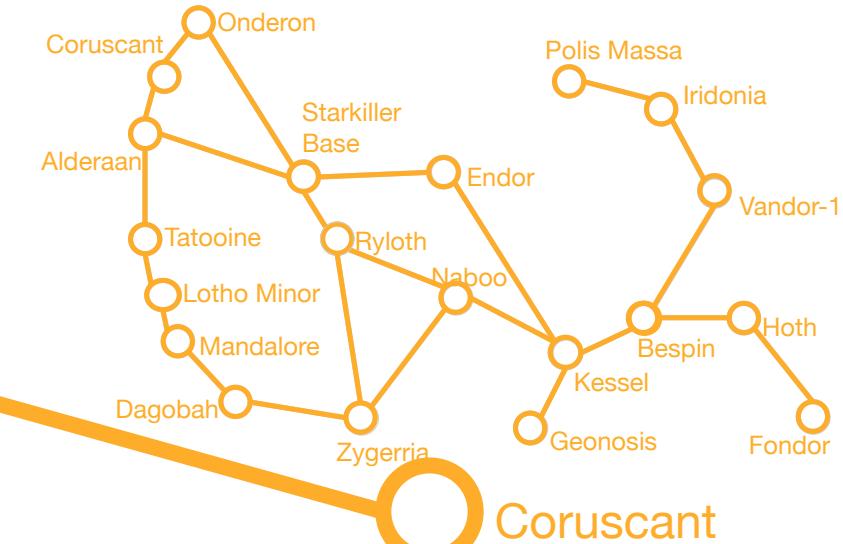
Greedy best-first search example

Initial State: Alderaan

Goal State: Kessel



Kessel	0
RylOTH	193
Starkiller Base	253
Tatooine	329
Alderaan	366
Coruscant	374
Onderon	380

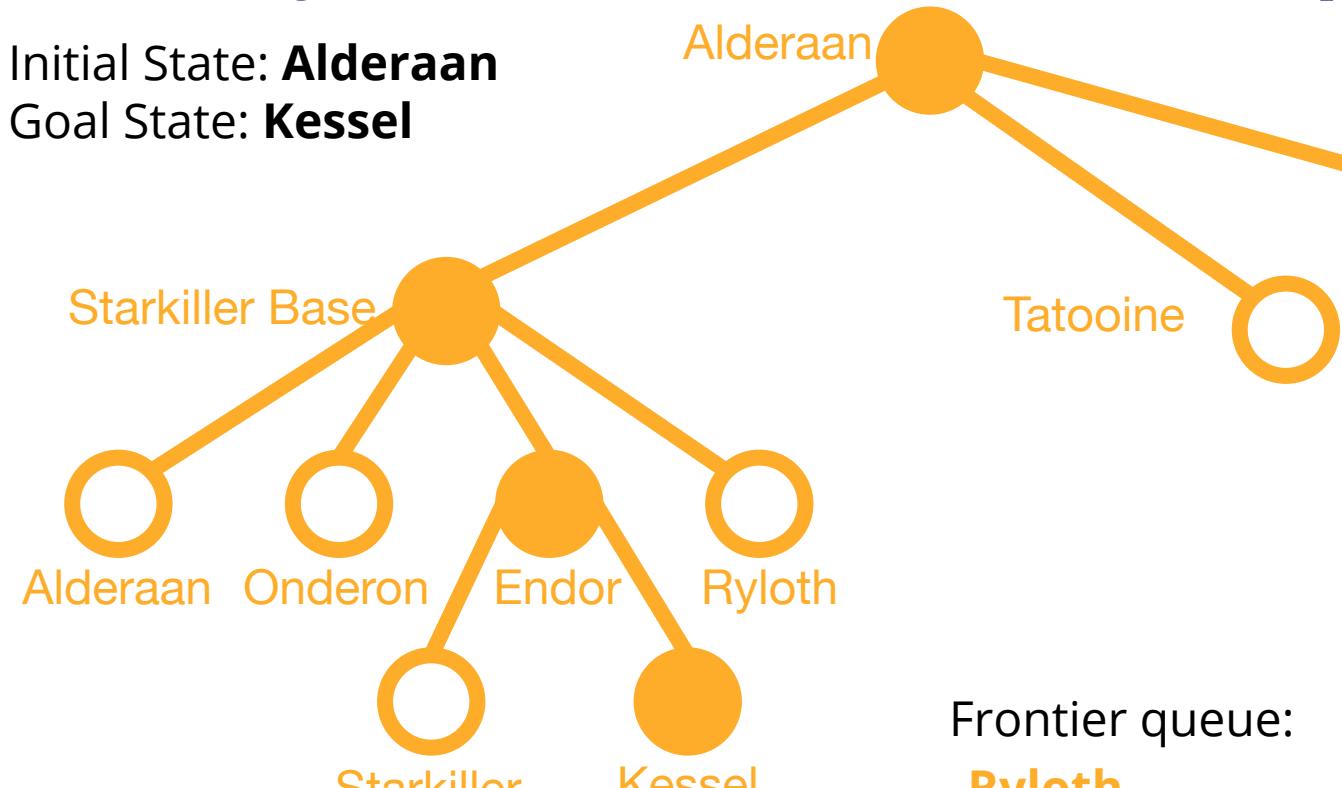


Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	RylOTH	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Greedy best-first search example

Initial State: Alderaan

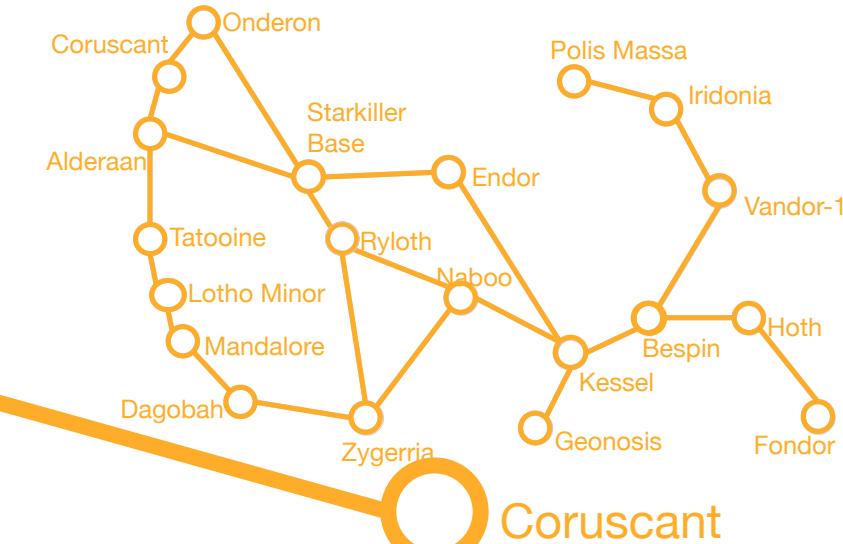
Goal State: Kessel



Goal reached !

Frontier queue:

RylOTH	193
Starkiller Base	253
Tatooine	329
Alderaan	366
Coruscant	374
Onderon	380

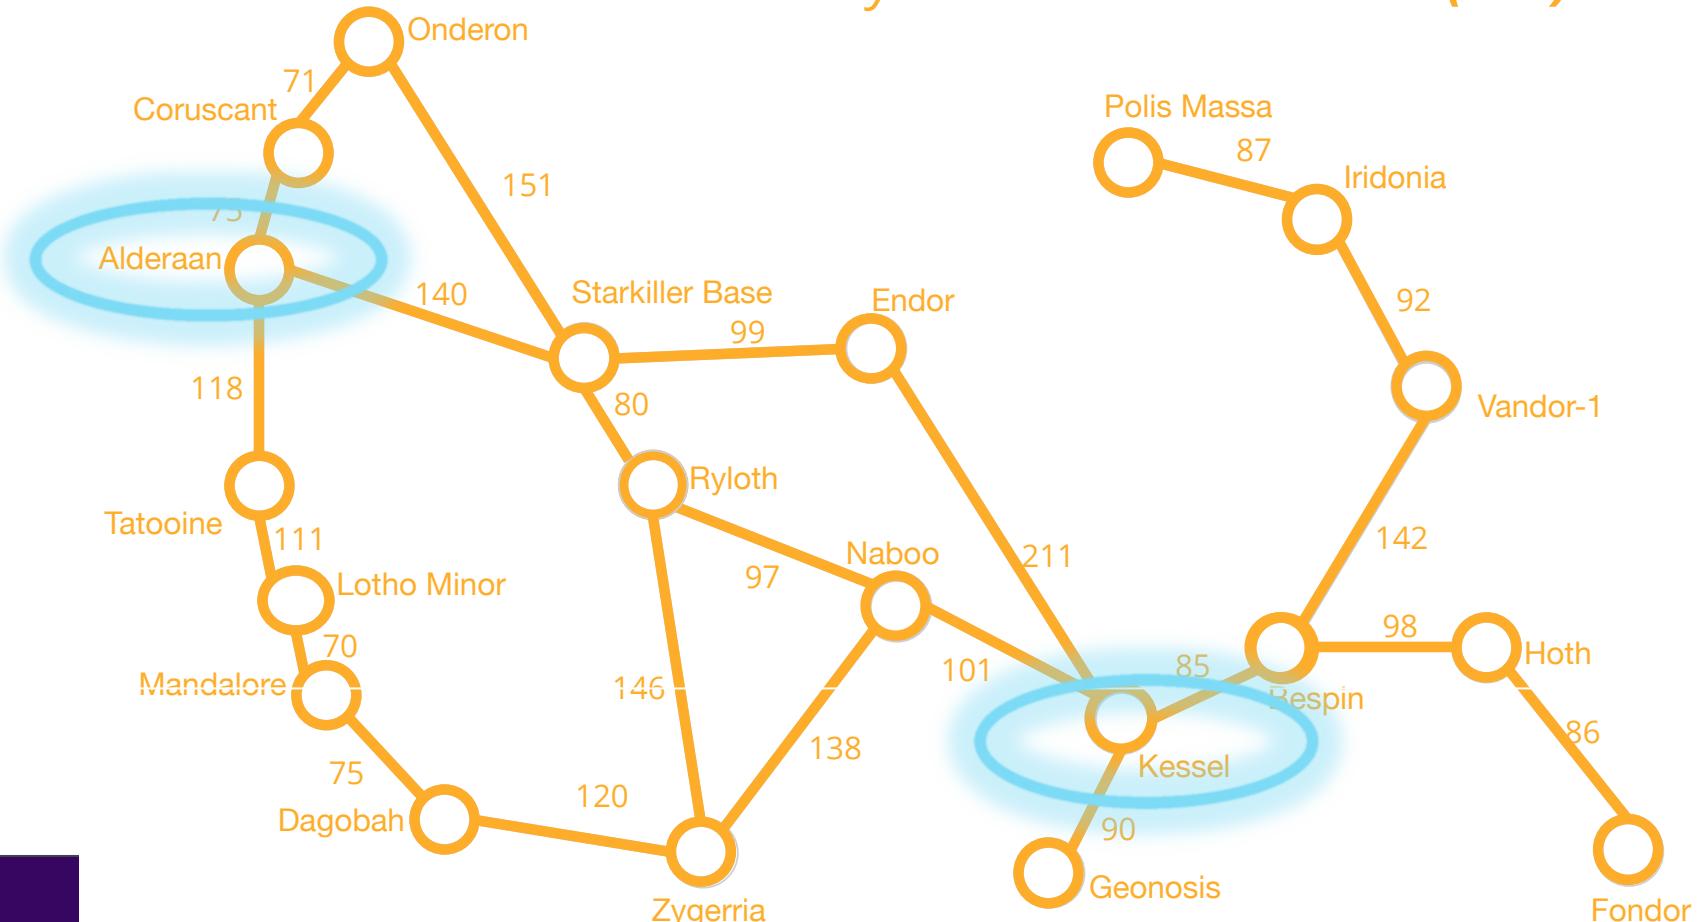


Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	RylOTH	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

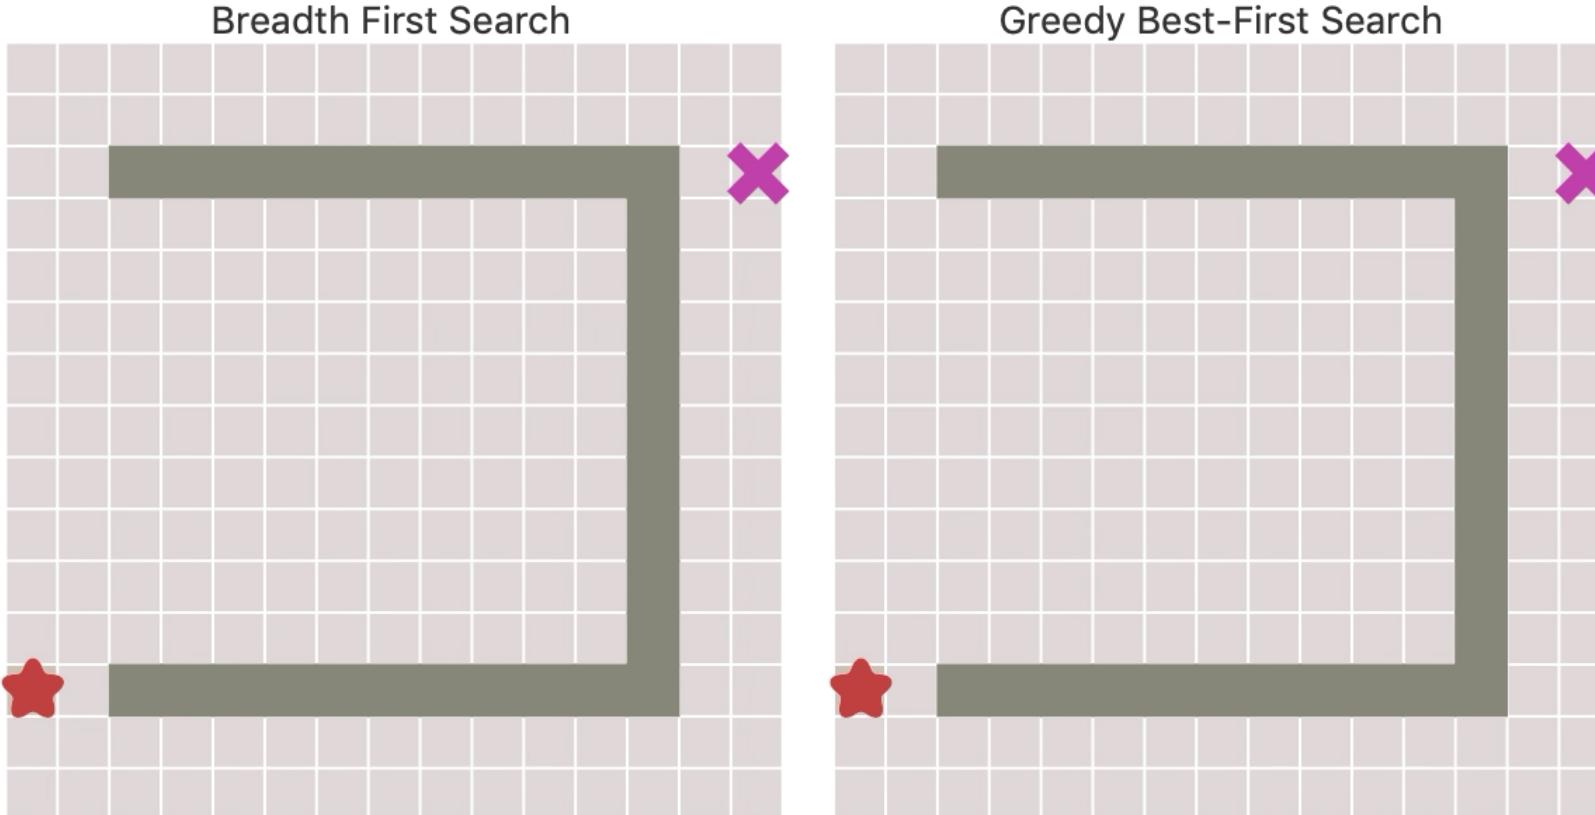
Properties of greedy best-first search

Optimal?

- No!
 - Found: *Alderaan → Starkiller Base → Endor → Kessel (450)*
 - Shorter: *Alderaan → Starkiller Base → Ryloth → Naboo → Kessel (418)*



BFS v. Greedy Best-First Search

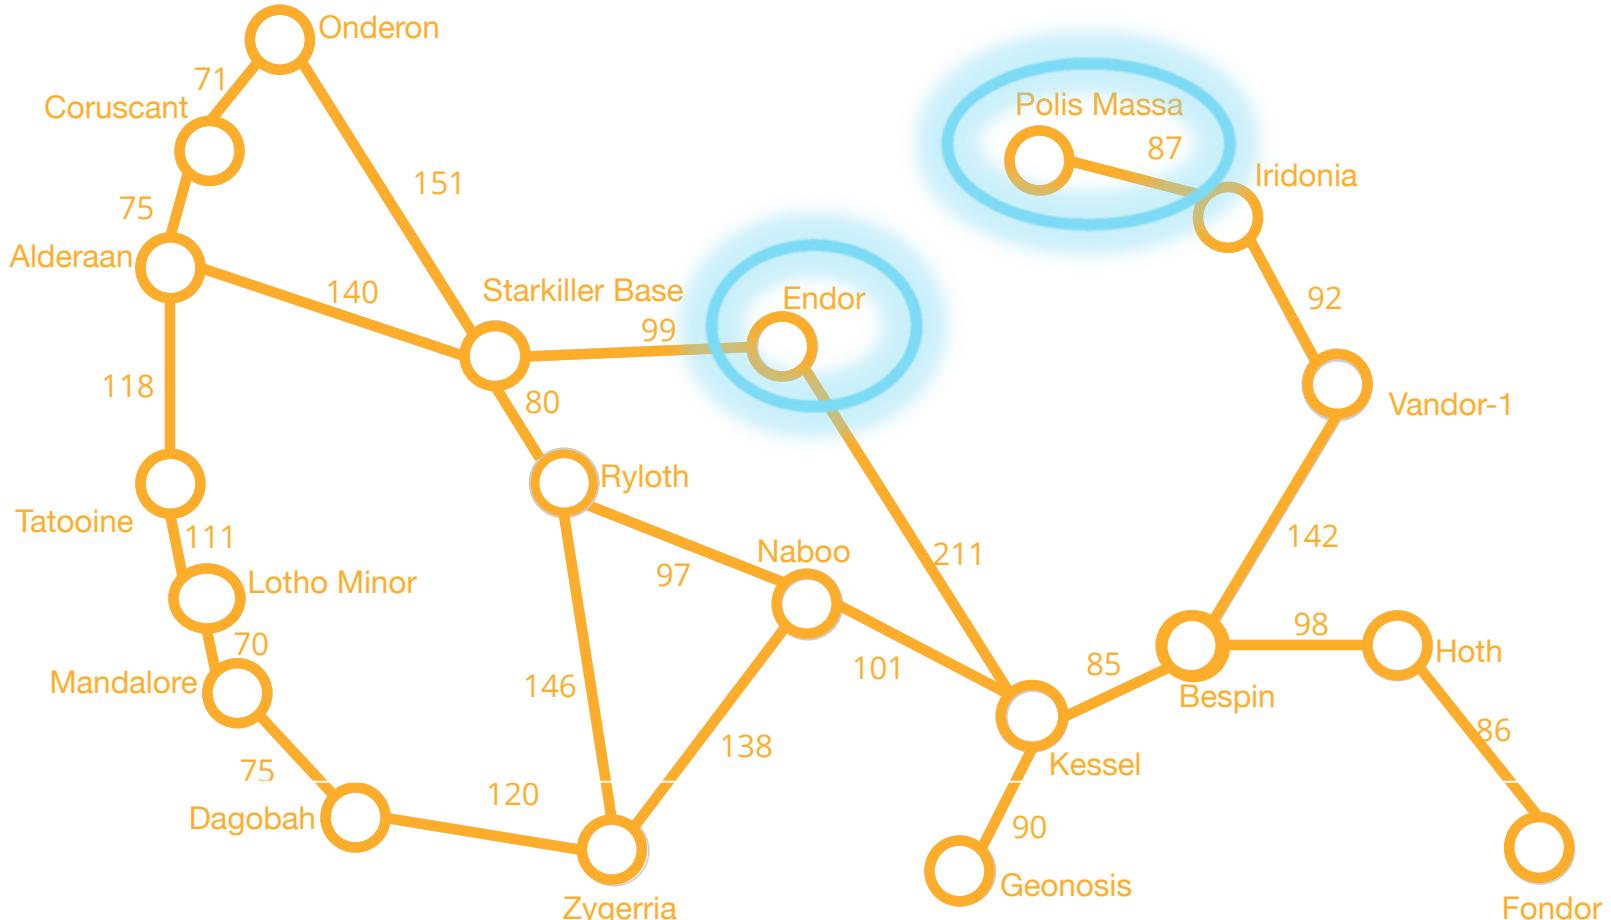


<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Properties of greedy best-first search

Complete?

- No – can get stuck in loops,
- e.g., Iridonia → Polis Massa → Iridonia → Polis Massa →...

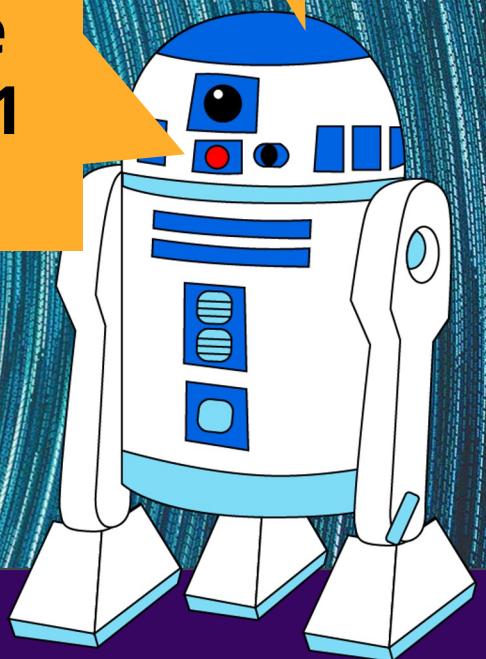


CIS 4210/5210:
ARTIFICIAL INTELLIGENCE

Informed Search

Quiz due on Sunday. HW3 is
due on Wednesday.

If you need an extension,
request is 24 hours in
advance via Canvas (no need
to email me). Usually, we
grant a max extension of 1
or 2 days.





A* search

AIMA 3.5



A* search

Best-known form of best-first search.

Key Idea: avoid expanding paths that are already expensive but expand most promising first.

Simple idea: $f(n) = g(n) + h(n)$

- $g(n)$ the actual cost (so far) to *reach* the node
- $h(n)$ estimated cost to *get from the node to the goal*
- $f(n)$ estimated *total cost* of path through n to goal

Implementation: Frontier queue as priority queue by increasing $f(n)$ (*as expected...*)

Key concept: Admissible heuristics

A heuristic $h(n)$ is *admissible* if it *never overestimates* the cost to reach the goal; i.e. it is *optimistic*

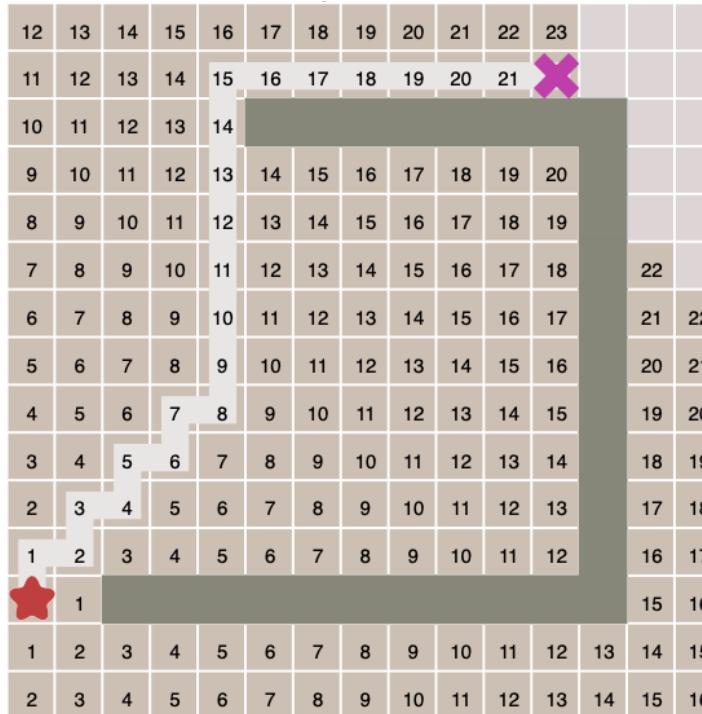
- Formally: $\forall n$, n a node:
 - $h(n) \leq h^*(n)$ where $h^*(n)$ is the true cost from n
 - $h(n) \geq 0$ so $h(G)=0$ for any goal G .

Example: $h_{SLD}(n)$ never overestimates the actual road distance

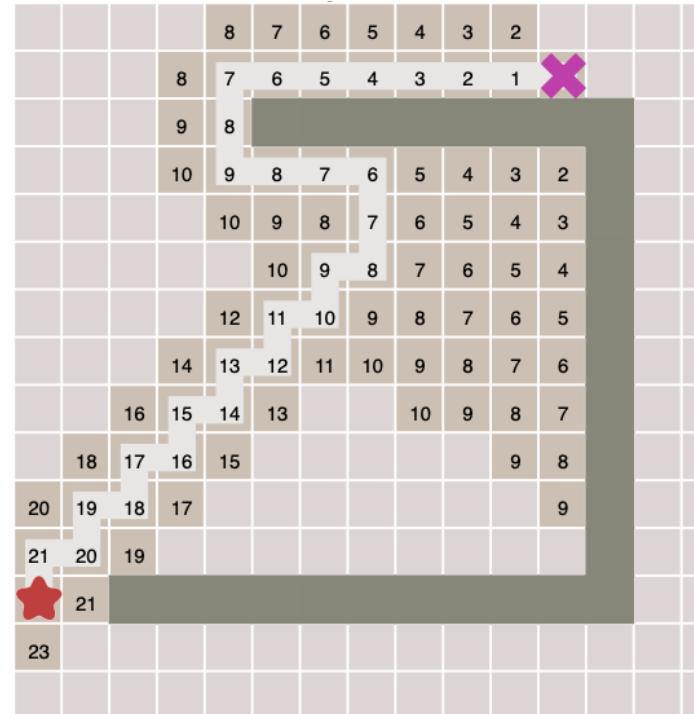
Theorem: If $h(n)$ is *admissible*, A* using **Tree Search** is *optimal*

A* is optimal with admissible heuristic

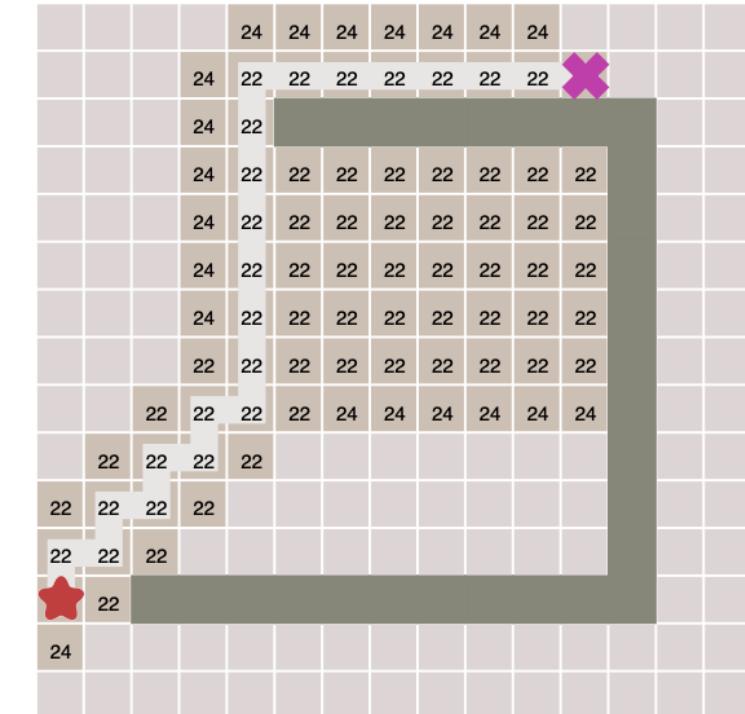
Uniform Cost Search



Greedy Best-First Search



A* Search



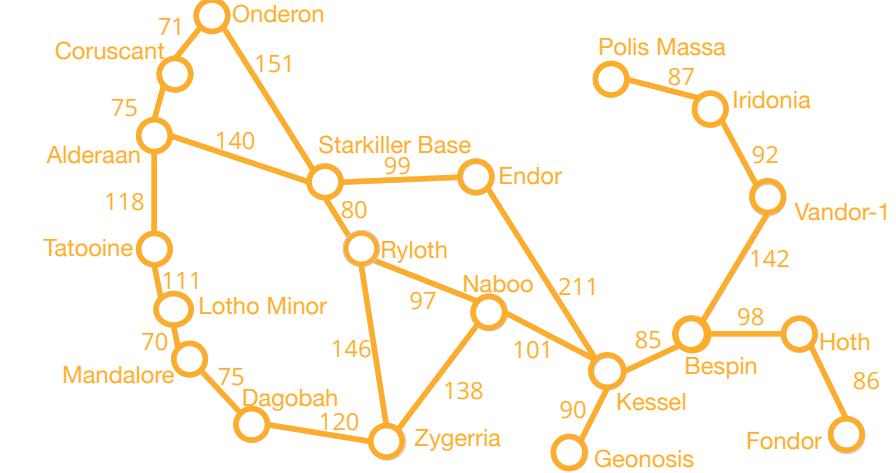
<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

A* search example

Initial State: Alderaan

Alderaan

Goal State: Kessel



Frontier queue:

Alderon $366 = 0 + 366$

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

A* search example

Initial State: **Alderaan**
Goal State: **Kessel**



$$f(n) = g(n) + h(n)$$

Frontier queue:

Starkiller Base $393 = 140 + 253$

Tatooine $447 = 118 + 329$

Coruscant $449 = 75 + 374$

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

A* search example

Initial State: **Alderaan**
Goal State: **Kessel**



Sort queue by $f(n)$.

$$f(n) = g(n) + h(n)$$

Frontier queue:

Starkiller Base $393 = 140 + 253$

Tatooine $447 = 118 + 329$

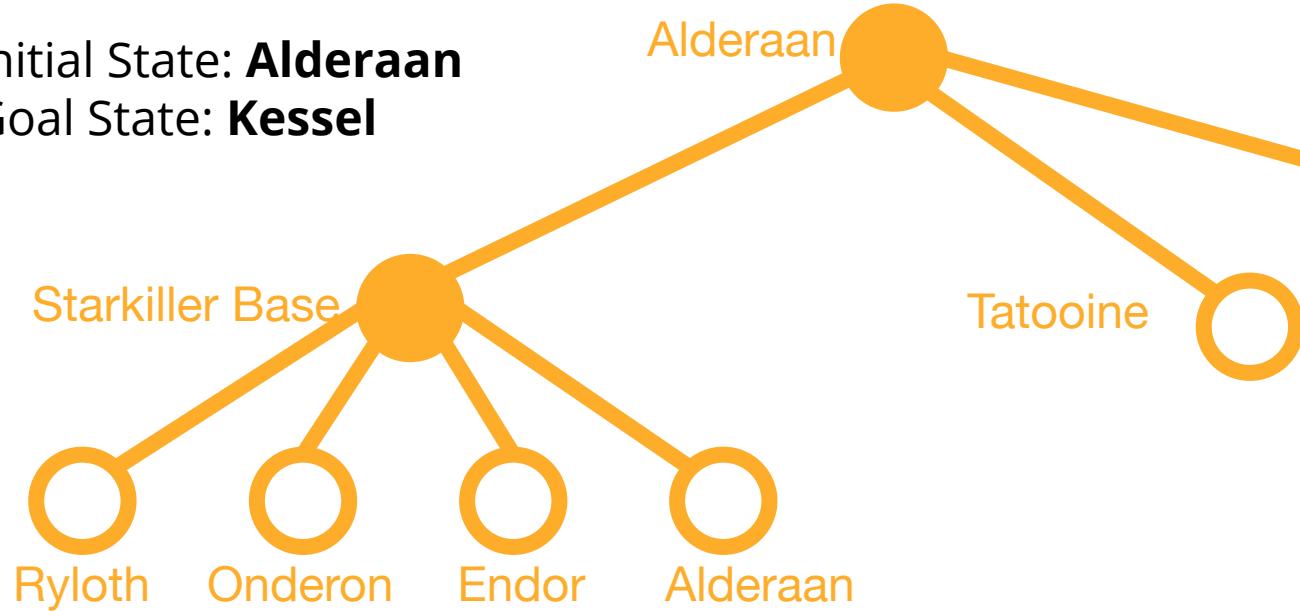
Coruscant $449 = 75 + 374$

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryleth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

A* search example

Initial State: Alderaan

Goal State: Kessel



Are we doing **graph** search or **tree** search?

Frontier queue:

Tatooine $447 = 118 + 329$

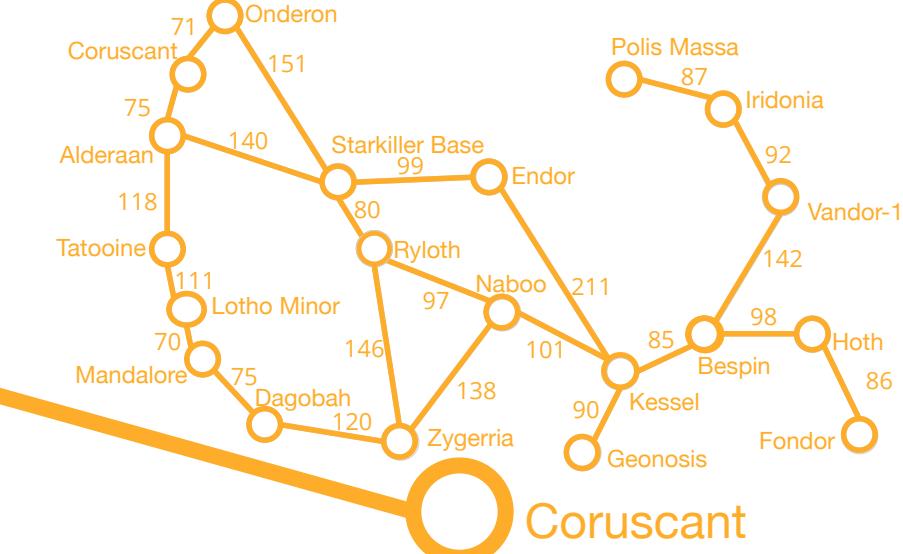
Coruscant $449 = 75 + 374$

Alderaan $646 = 280 + 366$

Onderon $526 = 146 + 380$

Endor $415 = 239 + 176$

RylOTH $413 = 220 + 193$



Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	RylOTH	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Graph search versus tree search

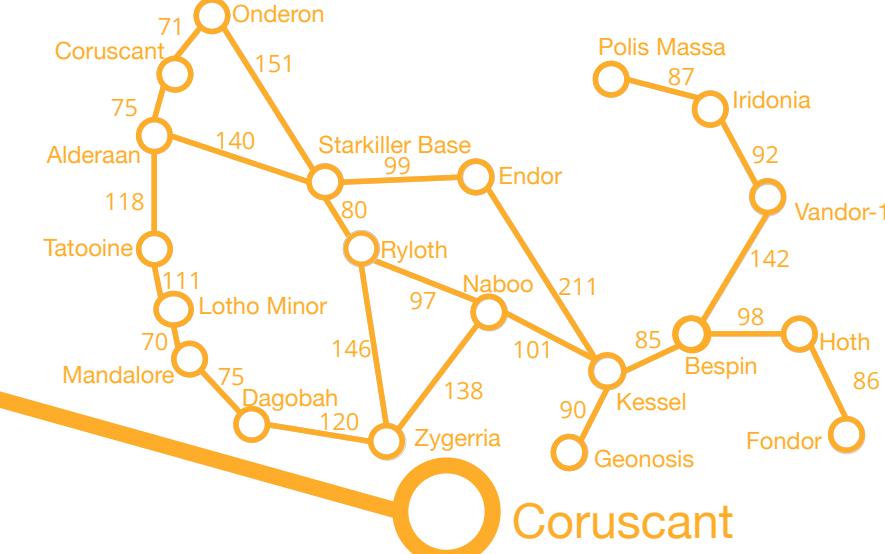
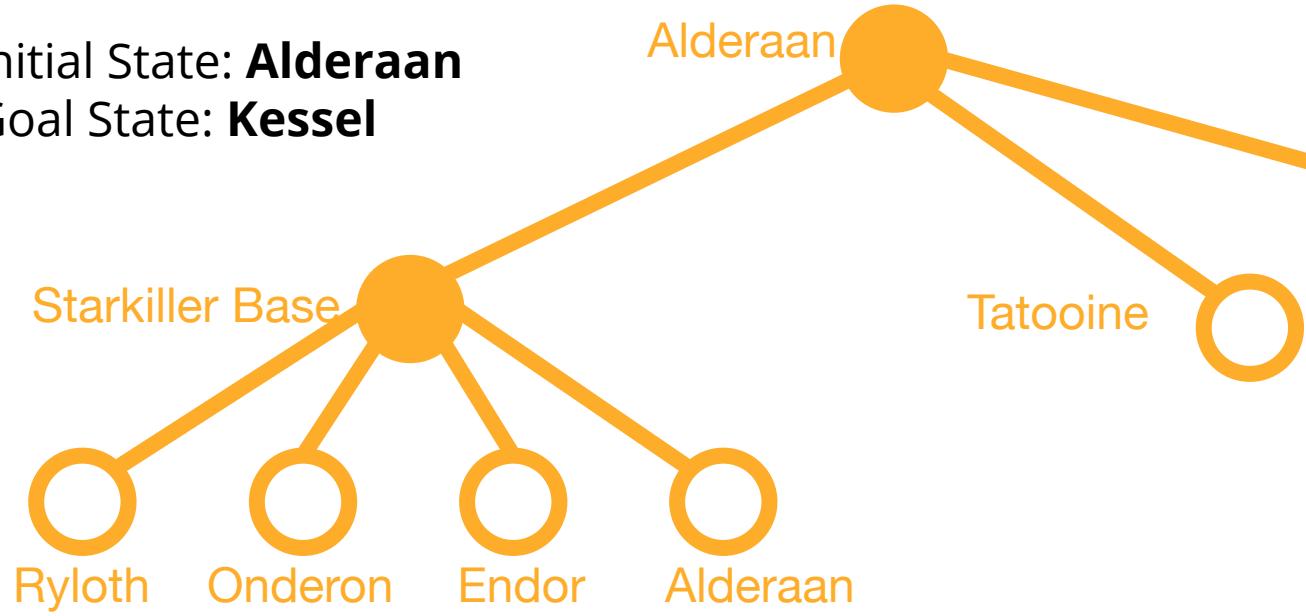


What is the difference between graph search and tree search?

PollEv.com/ccb

A* search example

Initial State: Alderaan
Goal State: Kessel



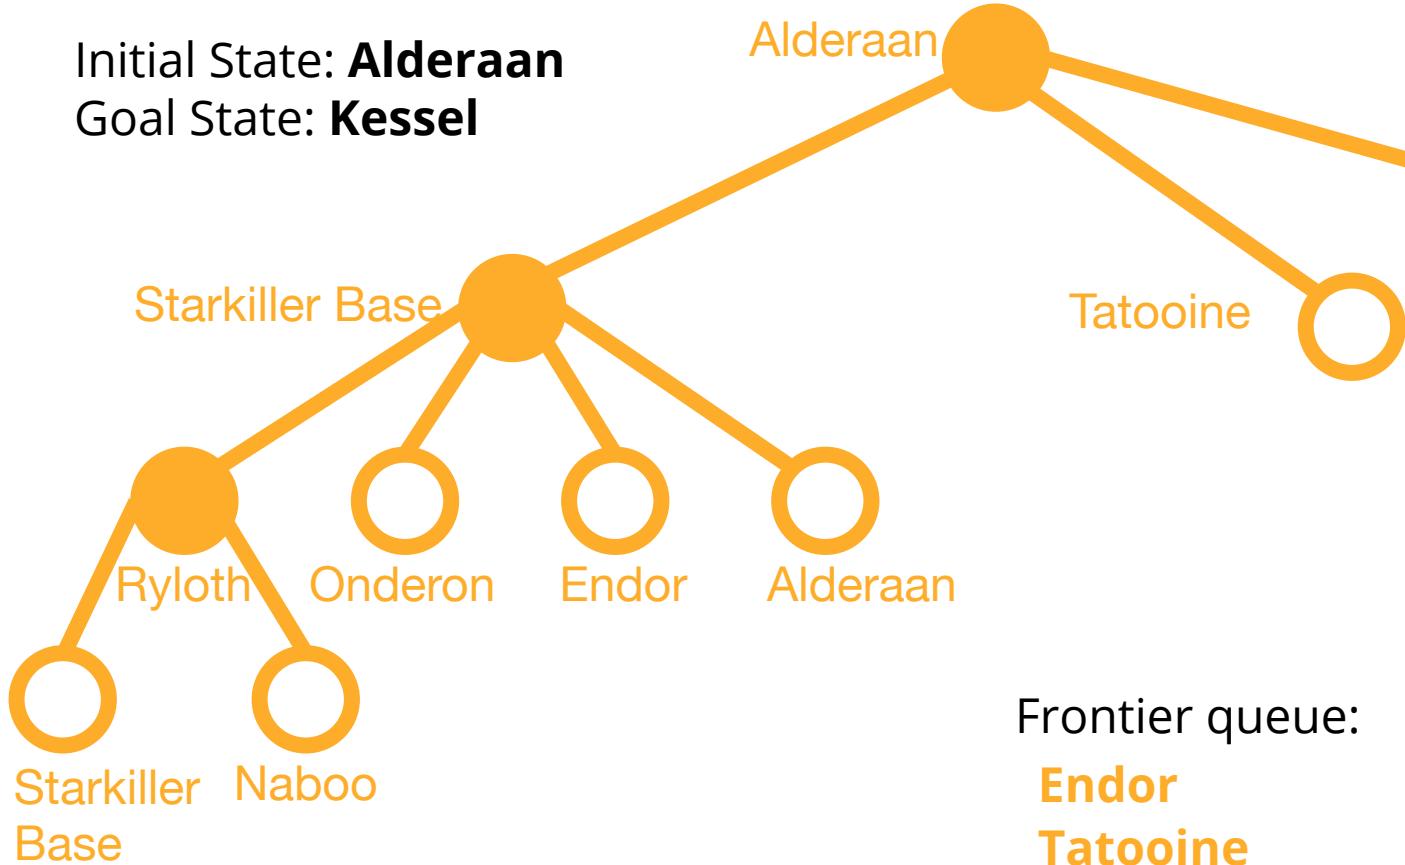
Frontier queue:

RylOTH	$413 = 220 + 193$
Endor	$415 = 239 + 176$
Tatooine	$447 = 118 + 329$
Coruscant	$449 = 75 + 374$
Onderon	$526 = 146 + 380$
Alderaan	$646 = 280 + 366$

Planet	SLD	Planet	SLD
Alderaan	366	Lotho Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	RylOTH	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

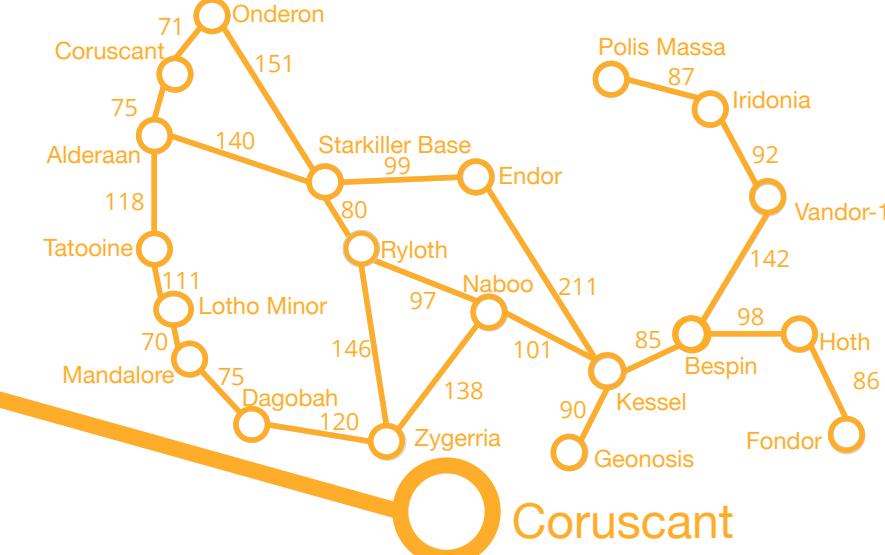
A* search example

Initial State: Alderaan
Goal State: Kessel



Frontier queue:

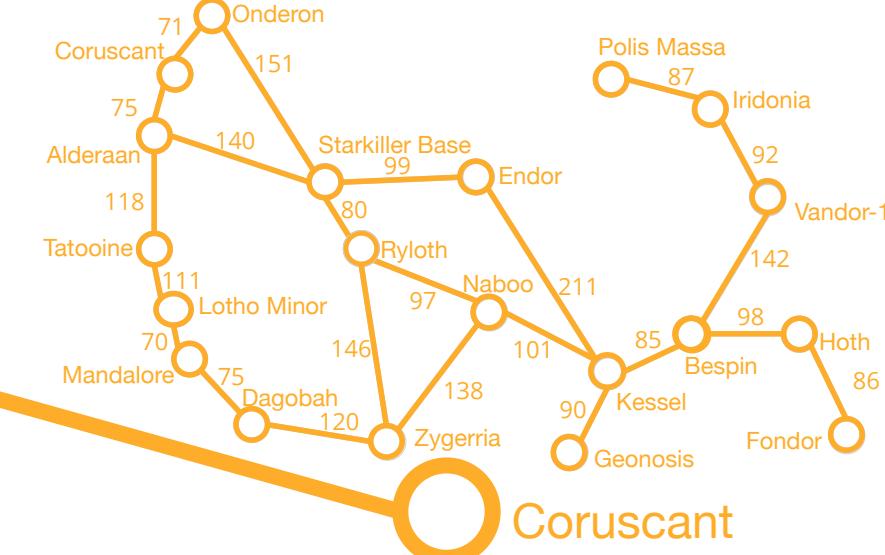
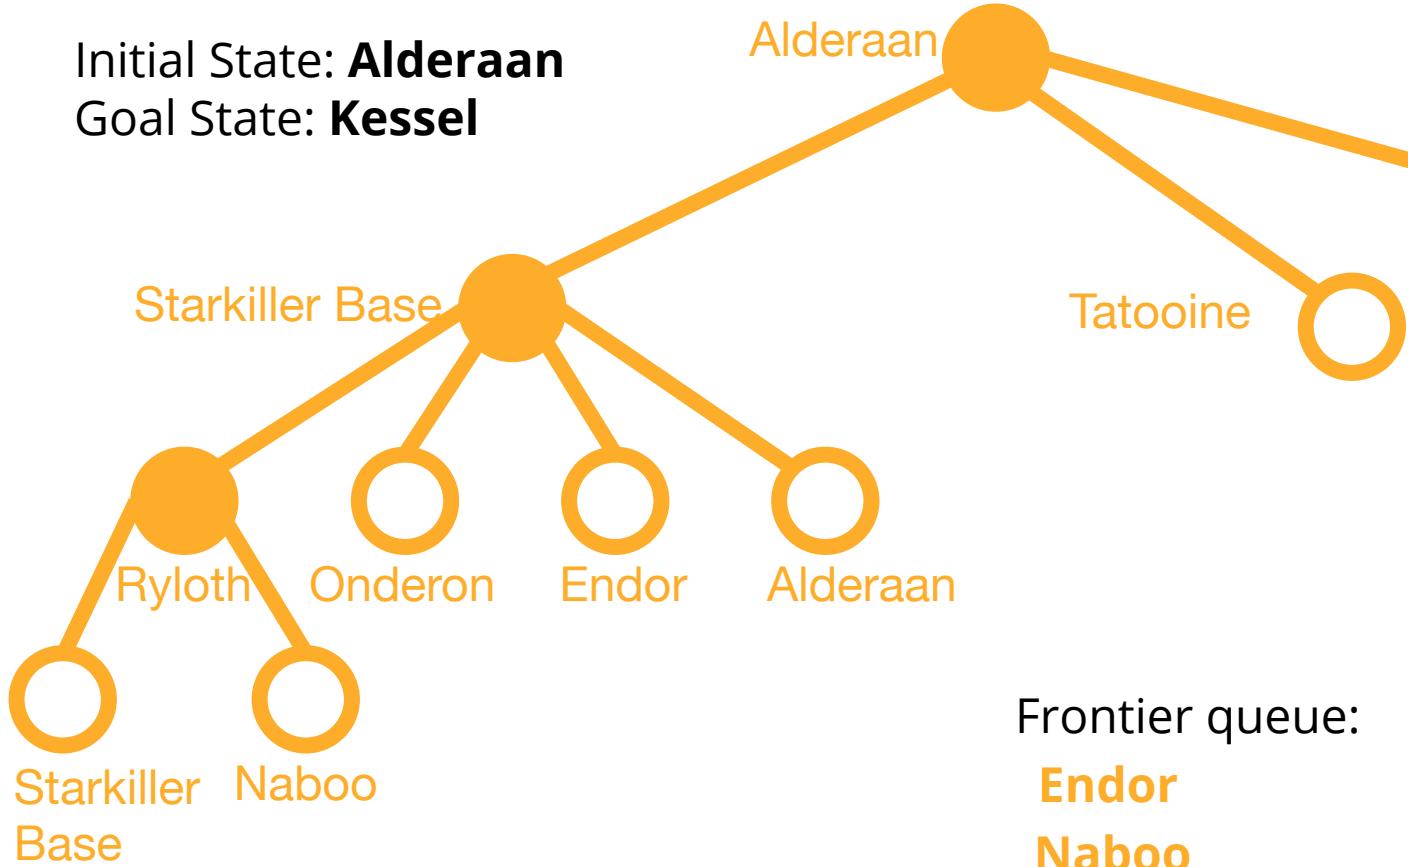
Alderaan	$366 = 280 + 366$
Starkiller Base	$532 = 300 + 253$
Tatooine	$447 = 118 + 329$
Coruscant	$449 = 75 + 374$
Onderon	$526 = 146 + 380$
Naboo	$417 = 317 + 100$



Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

A* search example

Initial State: Alderaan
Goal State: Kessel



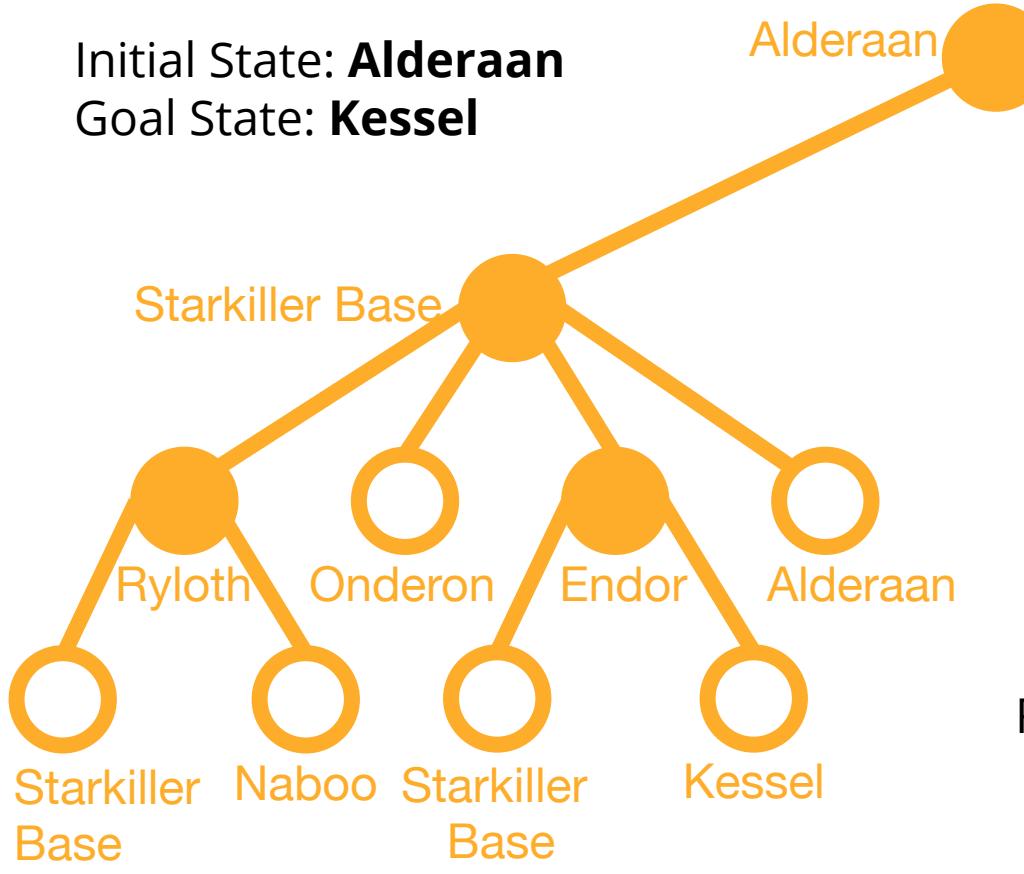
Frontier queue:

Alderaan	646 = 280 + 366
Starkiller Base	532 = 300 + 253
Onderon	526 = 146 + 380
Coruscant	449 = 75 + 374
Tatooine	447 = 118 + 329
Naboo	417 = 317 + 100
Rylloth	415 = 239 + 176

Planet	SLD	Planet	SLD
Alderaan	366	Lotho Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Rylloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

A* search example

Initial State: Alderaan
Goal State: Kessel



Do we apply the goal test to the **children** like in BFS?



No, we apply the goal test to the **top of the queue**.

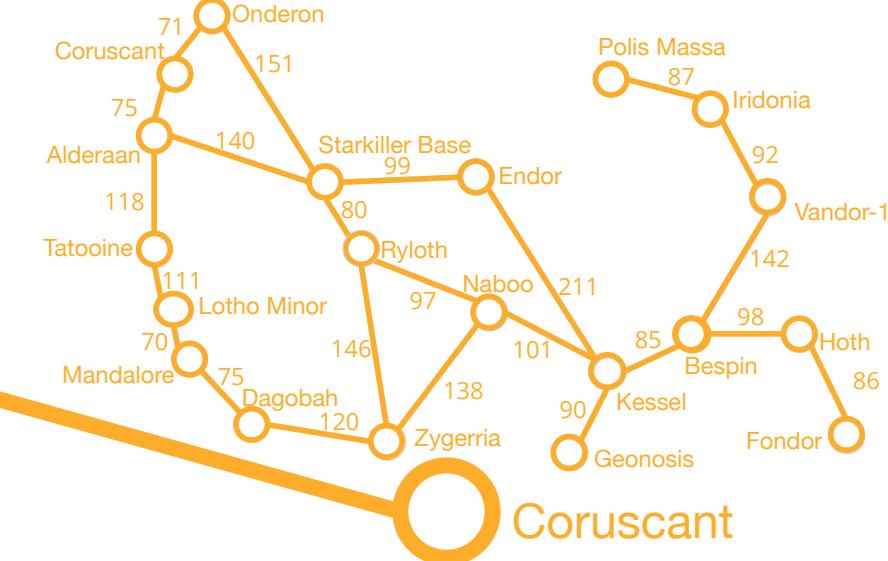
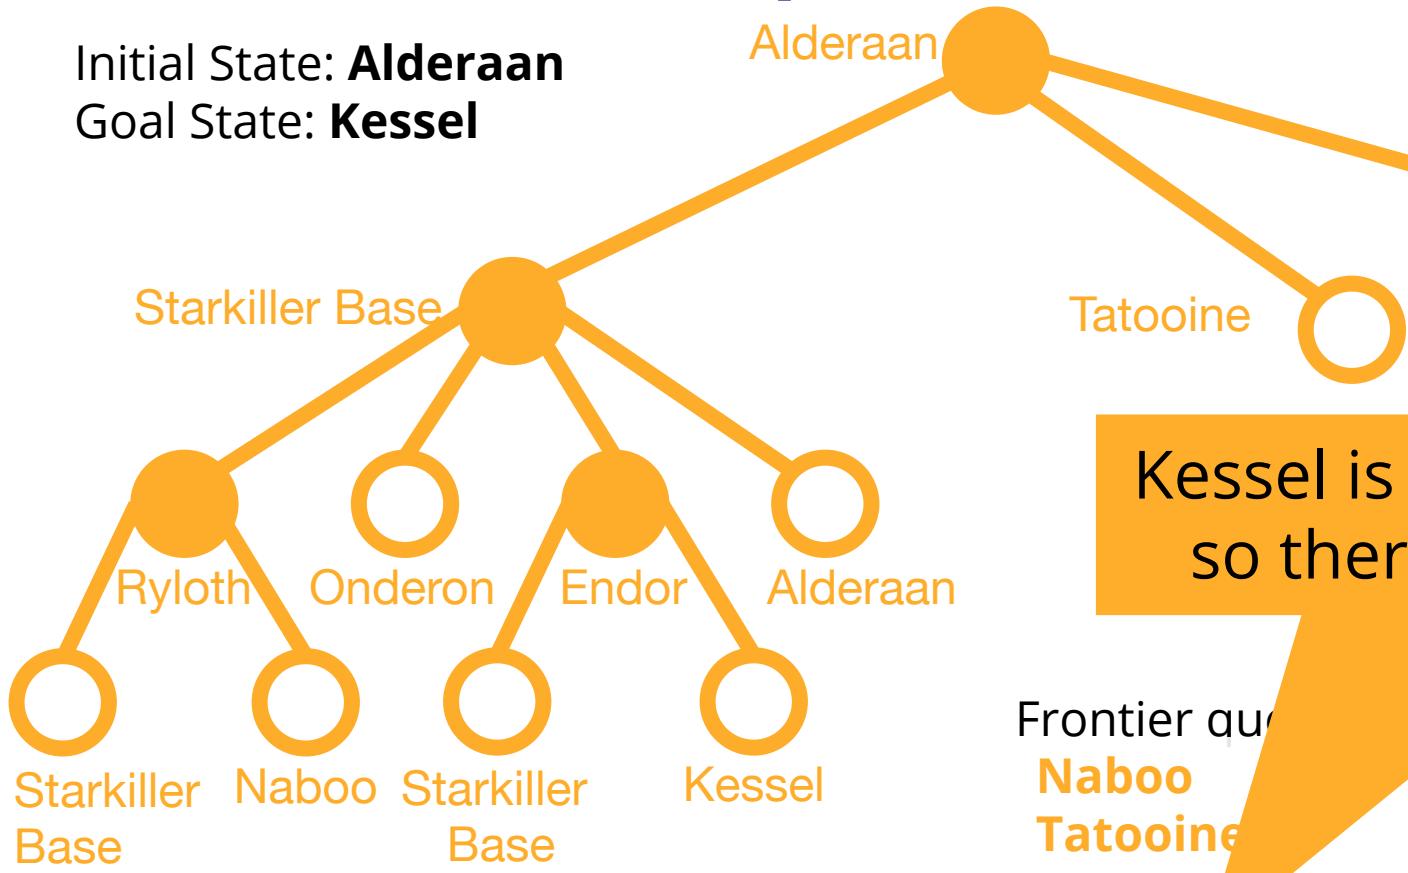
Frontier queue:

Naboo	$417 = 317 + 100$
Tatooine	$447 = 118 + 329$
Coruscant	$449 = 75 + 374$
Onderon	$526 = 146 + 380$
Starkiller Base	$532 = 300 + 253$
Alderaan	$646 = 280 + 366$
Starkiller Base	$591 = 338 + 253$
Kessel	$450 = 450 + 0$

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Rylloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

A* search example

Initial State: Alderaan
Goal State: Kessel



Kessel is not yet at **top of the queue**,
so there may be a **shorter path...**

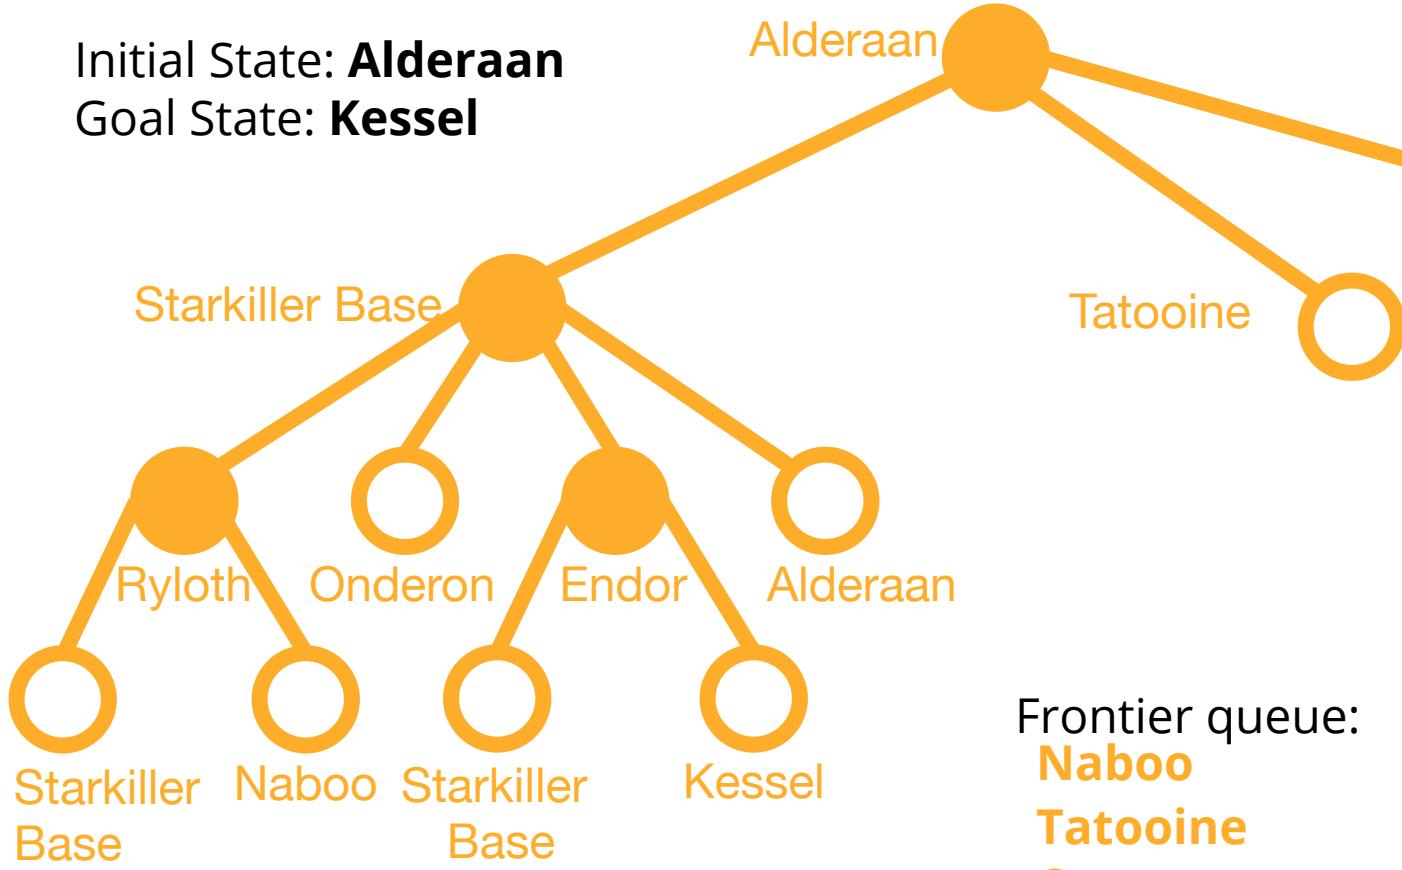
Frontier queue

Planet	SLD
Naboo	$417 = 317 + 100$
Tatooine	$447 = 118 + 329$
Coruscant	$449 = 75 + 374$
Kessel	$450 = 450 + 0$
Onderon	$526 = 146 + 380$
Starkiller Base	$532 = 300 + 253$
Alderaan	$646 = 280 + 366$
Starkiller Base	$591 = 338 + 253$

Planet	SLD
Alderaan	366
Bespin	80
Coruscant	374
Dagobah	242
Endor	176
Fondor	161
Geonosis	77
Hoth	151
Iridonia	226
Kessel	0
Lothal Minor	244
Mandalore	241
Naboo	100
Onderon	380
Polis Massa	234
Ryleth	193
Starkiller Base	253
Tatooine	329
Vandor-1	199
Zygerria	160

A* search example

Initial State: **Alderaan**
 Goal State: **Kessel**



Frontier queue:

Naboo

$$417 = 317 + 100$$

Tatooine

$$447 = 118 + 329$$

Coruscant

$$449 = 75 + 374$$

Kessel

$$450 = 450 + 0$$

Onderon

$$526 = 146 + 380$$

Starkiller Base

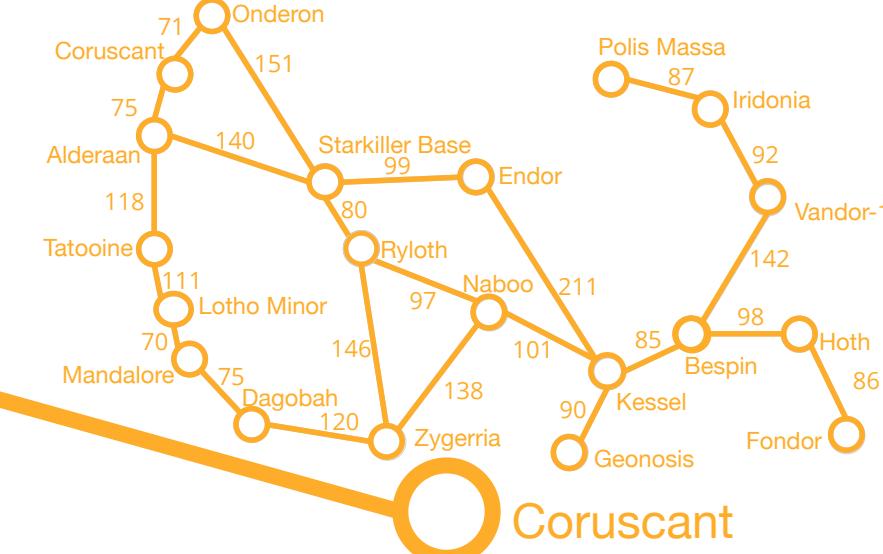
$$532 = 300 + 253$$

Starkiller Base

$$591 = 338 + 253$$

Alderaan

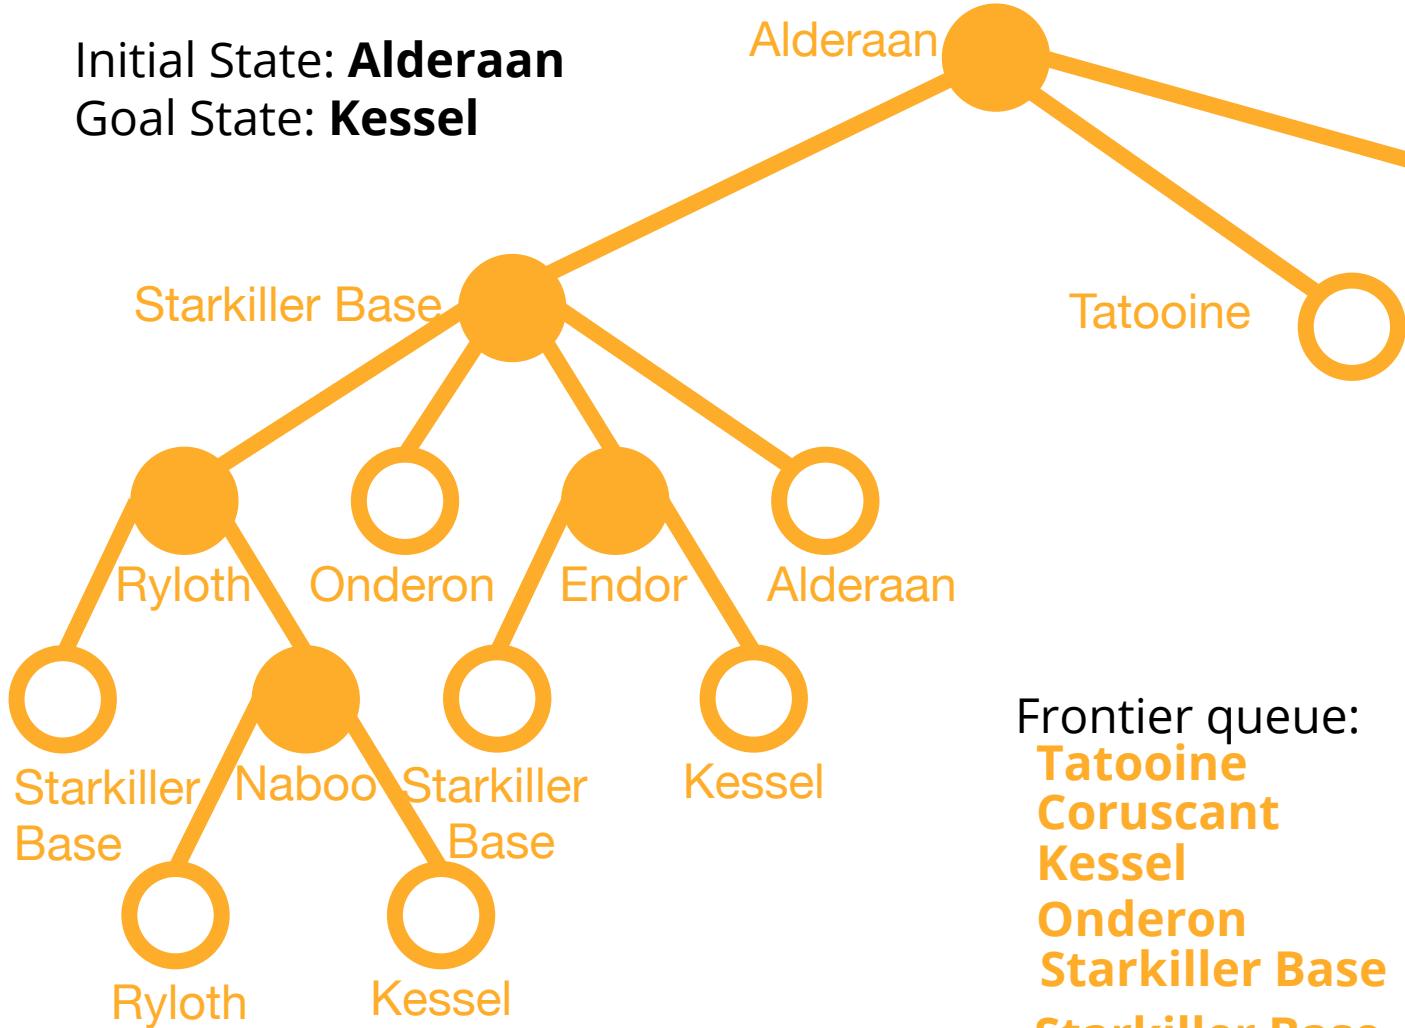
$$646 = 280 + 366$$



Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Ryloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

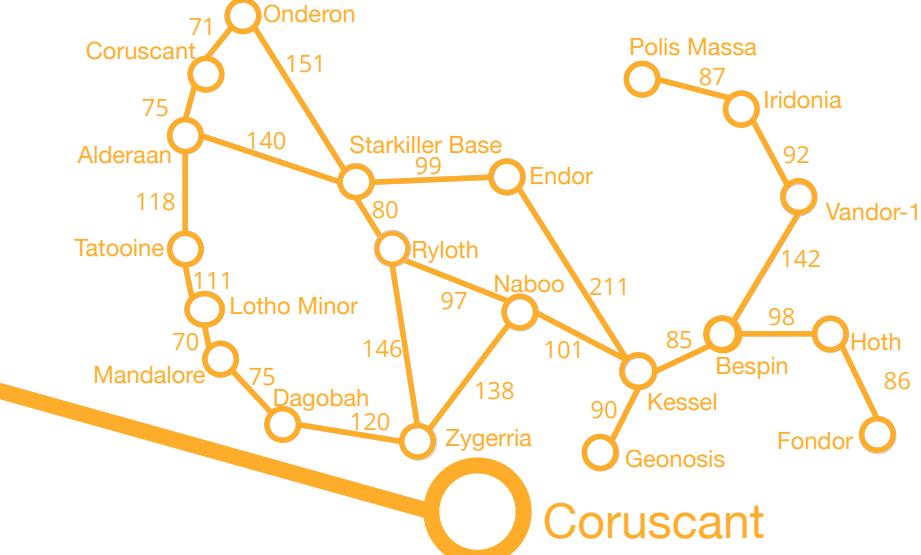
A* search example

Initial State: **Alderaan**
 Goal State: **Kessel**



Frontier queue:

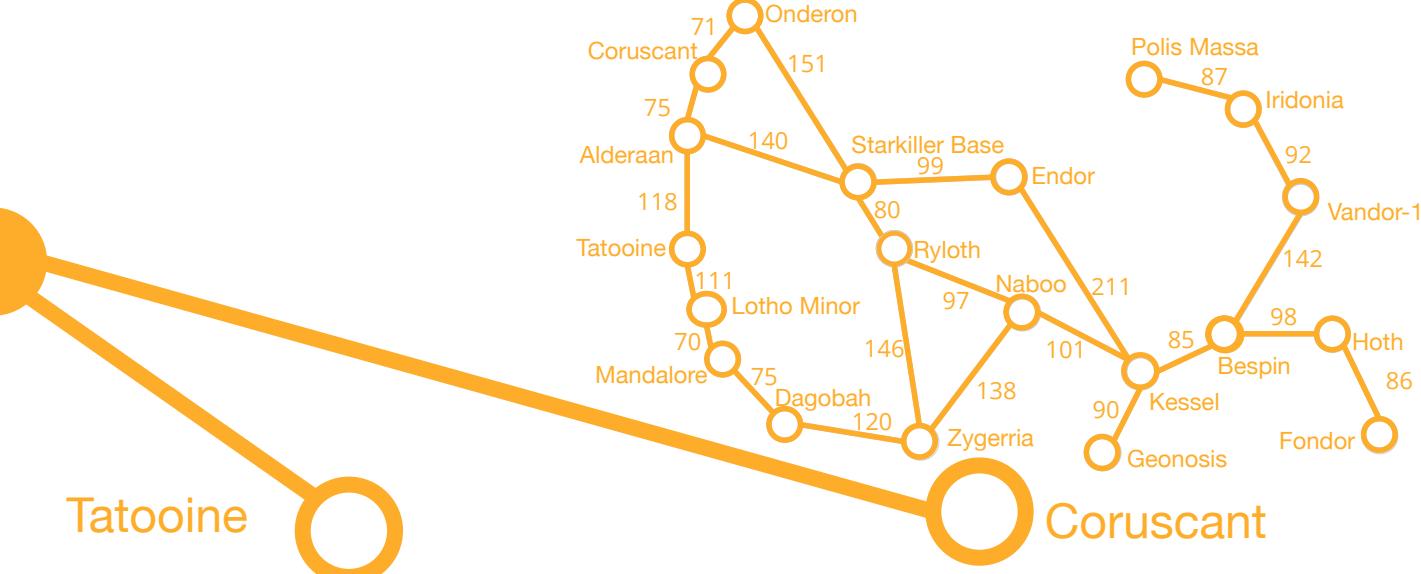
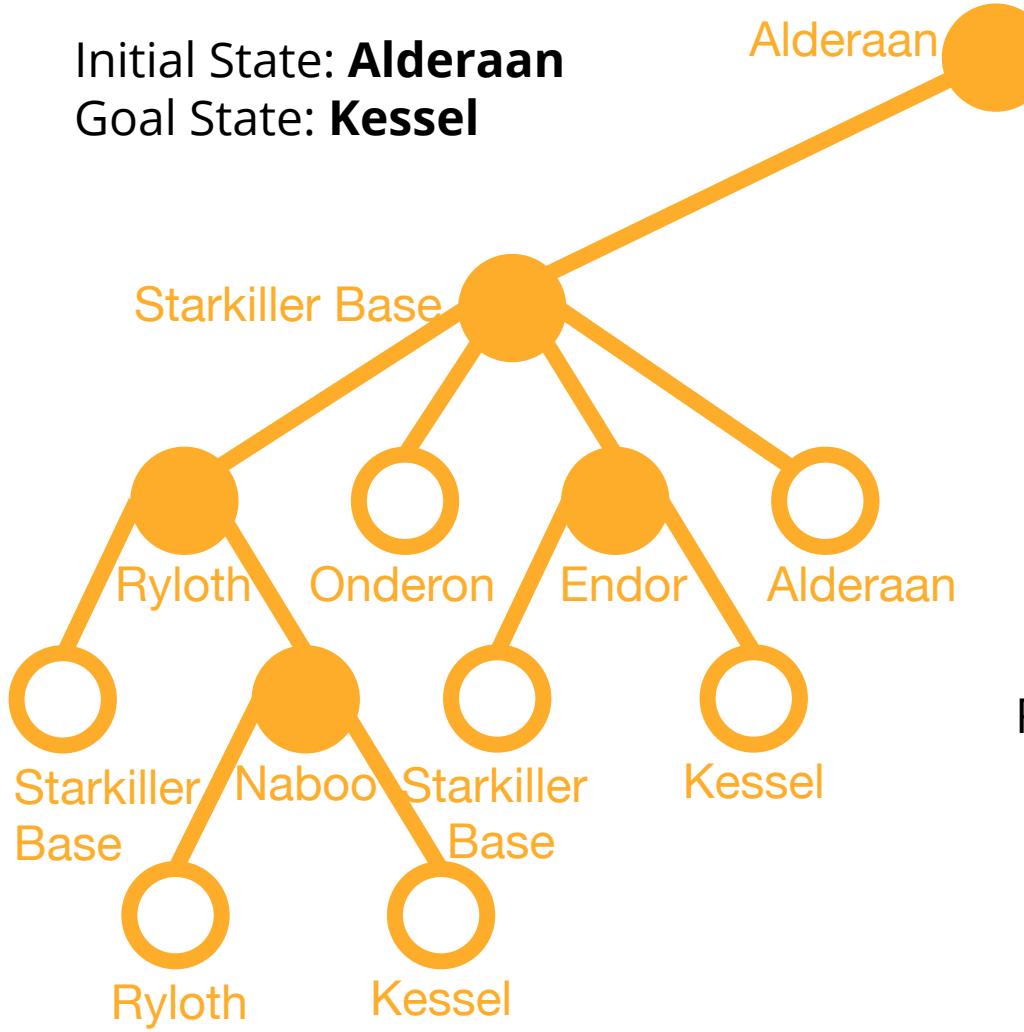
Tatooine
Coruscant
Kessel
Onderon
Starkiller Base
Starkiller Base
Alderaan
Rylloth
Kessel



447 = 118 + 329
449 = 75 + 374
450 = 450 + 0
526 = 146 + 380
532 = 300 + 253
591 = 338 + 253
646 = 280 + 366
607 = 414 + 193
418 = 418 + 0

A* search example

Initial State: Alderaan
Goal State: Kessel



Passes the goal test, so **this** is the shortest path to Kessel!

Frontier queue...

Planet	SLD
Kessel	$418 = 418 + 0$
Tatooine	$447 = 118 + 329$
Coruscant	$449 = 75 + 374$
Kessel	$450 = 450 + 0$
Onderon	$526 = 146 + 380$
Starkiller Base	$532 = 300 + 253$
Starkiller Base	$591 = 338 + 253$
Rylloth	$607 = 414 + 193$
Alderaan	$646 = 280 + 366$

Planet	SLD	Planet	SLD
Alderaan	366	Lothal Minor	244
Bespin	80	Mandalore	241
Coruscant	374	Naboo	100
Dagobah	242	Onderon	380
Endor	176	Polis Massa	234
Fondor	161	Rylloth	193
Geonosis	77	Starkiller Base	253
Hoth	151	Tatooine	329
Iridonia	226	Vandor-1	199
Kessel	0	Zygerria	160

Idea: Admissibility



Inadmissible
(pessimistic) heuristics
break optimality by
trapping good plans on
the frontier



Admissible (optimistic)
heuristics slow down
bad plans but never
outweigh true costs

Heuristic functions

For the 8-puzzle

- **Avg. solution cost is about 22 steps**
 - (branching factor ≤ 3)
 - (branching factor ≤ 3)
 - A good heuristic function can reduce the search process



Start State



Goal State

Example Admissible heuristics

For the 8-puzzle:

$h_{oop}(n)$ = number of out of place tiles

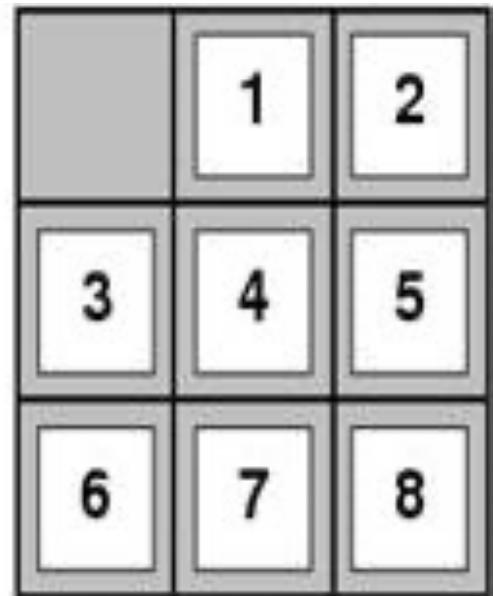
$h_{md}(n)$ = total Manhattan distance (i.e., #
of moves from desired location of
each tile)

$$h_{oop}(S) = 8$$

$$h_{md}(S) = 3+1+2+2+2+3+3+2 = 18$$



Start State



Goal State

Relaxed problems

A problem with fewer restrictions on the actions than the original is called a *relaxed problem*

The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem

If the rules of the 8-puzzle are relaxed so that a tile can move *anywhere*, then $h_{oop}(n)$ gives the shortest solution

If the rules are relaxed so that a tile can move to *any adjacent square*, then $h_{md}(n)$ gives the shortest solution

Defining Heuristics: $h(n)$

Cost of an exact solution to a *relaxed* problem (fewer restrictions on operator)

Constraints on *Full* Problem:

A tile can move from square A to square B *if A is adjacent to B and B is blank.*

- Constraints on *relaxed* problems:
 - A tile can move from square A to square B *if A is adjacent to B.* (h_{md})
 - A tile can move from square A to square B *if B is blank.*
 - A tile can move from square A to square B. (h_{oop})

Dominance: A metric on better heuristics

If $h_2(n) \geq h_1(n)$ for all n (both admissible)

- then h_2 *dominates* h_1

So h_2 is optimistic, but more accurate than h_1

- h_2 is therefore better for search
- Notice: h_{md} dominates h_{oop}

Typical search costs (average number of nodes expanded):

$d=12$ Iterative Deepening Search = 3,644,035 nodes

$A^*(h_{oop}) = 227$ nodes, $A^*(h_{md}) = 73$ nodes

$d=24$ IDS = too many nodes

$A^*(h_{oop}) = 39,135$ nodes, $A^*(h_{md}) = 1,641$ nodes

The best and worst admissible heuristics

$h^*(n)$ - the (unachievable) Oracle heuristic

- $h^*(n)$ = the true distance from the n to goal

$h_{\text{we're here already}}(n) = h_{\text{teleportation}}(n) = 0$

Admissible: both yes!!!

$h^*(n)$ *dominates all other heuristics*

$h_{\text{teleportation}}(n)$ *is dominated by all heuristics*

A* search is Optimal

AIMA 3.5



A* search

Best-known form of best-first search.

Key Idea: avoid expanding paths that are already expensive, but expand most promising first.

Simple idea: $f(n) = g(n) + h(n)$

- $g(n)$ the actual cost (so far) to *reach* the node
- $h(n)$ estimated cost to *get from the node to the goal*
- $f(n)$ estimated *total cost* of path through n to goal

Implementation: Frontier queue as priority queue by increasing $f(n)$ (*as expected...*)

Key: Admissibility



Inadmissible (pessimistic) heuristics break optimality by pushing good plans too far back on the frontier, which means they may never get expanded.



Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs. That means that the true best plan will always be expanded.

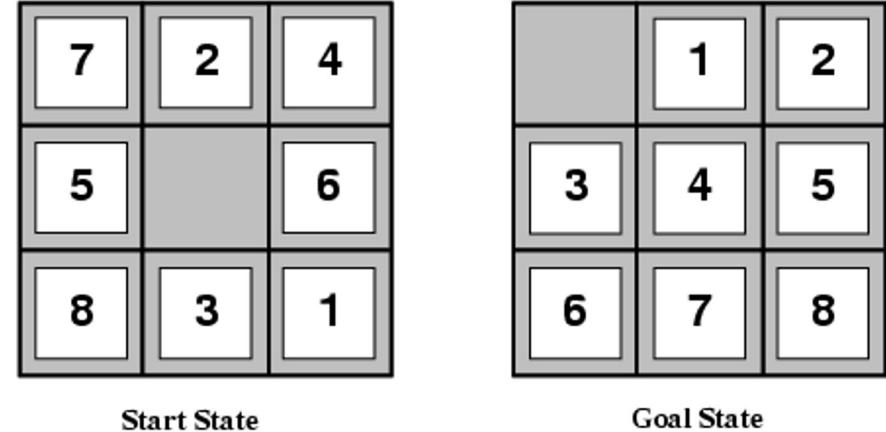
Admissible Heuristics

A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

Is Manhattan Distance admissible?



Coming up with admissible heuristics is most of what's involved in using A* in practice.

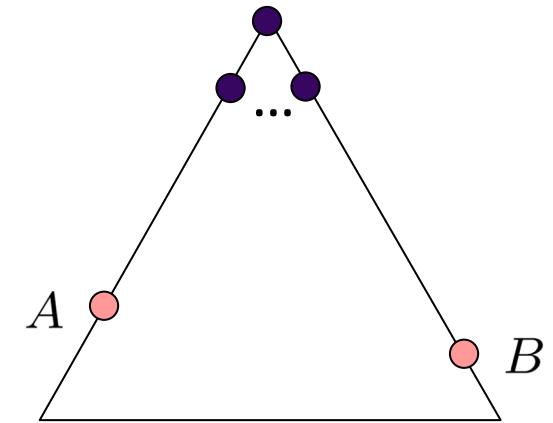
Optimality of A* Tree Search

Assume:

A is an optimal goal node

B is a suboptimal goal node

h is admissible



Claim:

A will exit the frontier before B

Slide credit: Dan Klein and Pieter

Abbeel

<http://ai.berkeley.edu>

Optimality of A* Tree Search

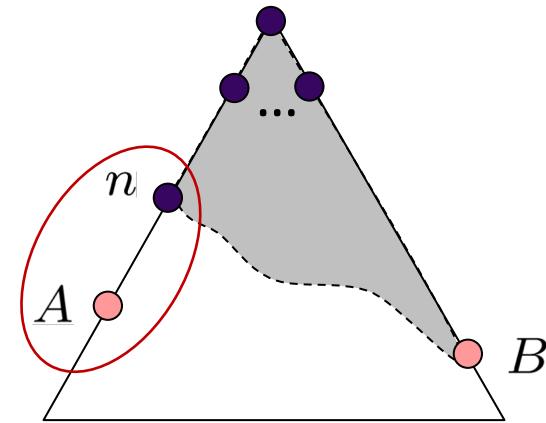
Proof:

Imagine B is on the frontier

Some ancestor n of A is on the frontier, too (maybe A!)

Claim: n will be expanded before B

- $f(n)$ is less or equal to $f(A)$



$$f(n) = g(n) + h(n) \quad \text{Definition of f-cost}$$

$$f(n) \leq g(A) \quad \text{Admissibility of } h$$

$$g(A) = f(A) \quad h = 0 \text{ at a goal}$$

Slide credit: Dan Klein and Pieter

Abbeel

<http://ai.berkeley.edu>

Optimality of A* Tree Search

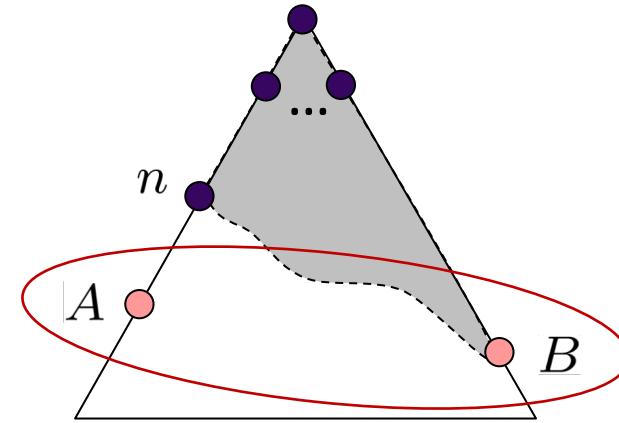
Proof:

Imagine B is on the frontier

Some ancestor n of A is on the frontier, too (maybe A!)

Claim: n will be expanded before B

- $f(n)$ is less or equal to $f(A)$
- $f(A)$ is less than $f(B)$



$g(A) < g(B)$ B is suboptimal
 $f(A) < f(B)$ $h = 0$ at a goal

Slide credit: Dan Klein and Pieter

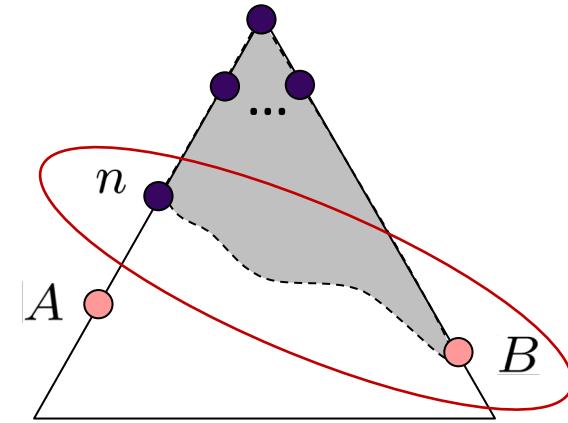
Abbeel

<http://ai.berkeley.edu>

Optimality of A* Tree Search

Proof:

- Imagine B is on the frontier
- Some ancestor n of A is on the frontier, too (maybe A!)
- Claim: n will be expanded before B
 - $f(n)$ is less or equal to $f(A)$
 - $f(A)$ is less than $f(B)$
 - n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal



$$f(n) \leq f(A) < f(B)$$

Slide credit: Dan Klein and Pieter

Abbeel

<http://ai.berkeley.edu>

Consistent Heuristics

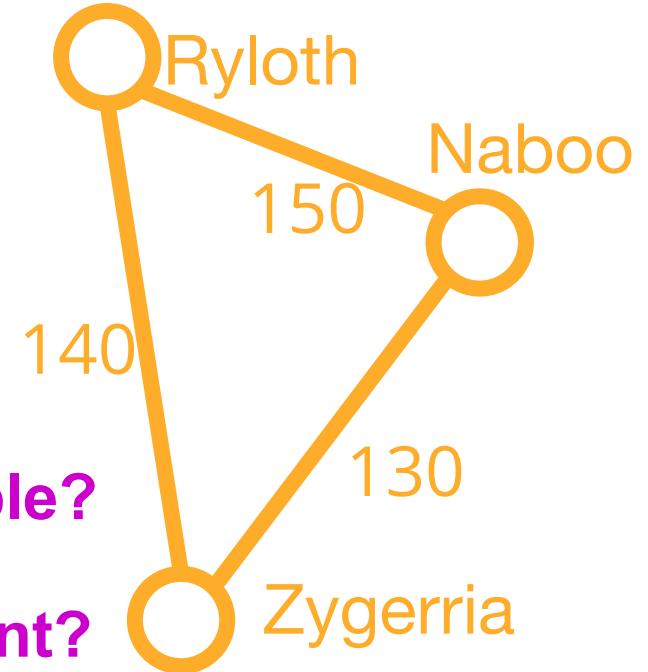
This ensures that A* is optimal in **graph search** in addition to tree search.

Consistent heuristics are a subset of admissible heuristics, which provide even stronger guarantees.

They obey the **triangle inequality**.

Planet	h_1
Ryloth	150
Zygerria	5
Naboo	0

Planet	h_2
Ryloth	150
Zygerria	120
Naboo	0



Is h_1 admissible?

Is h_1 consistent?

Is h_2 admissible?

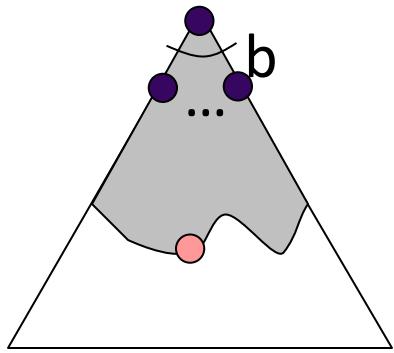
Is h_2 consistent?

Properties of A*

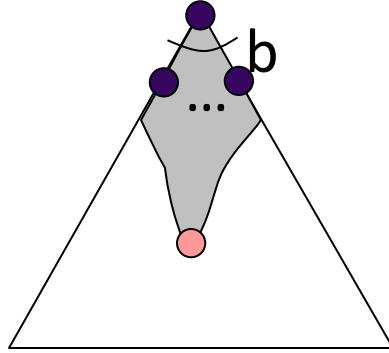
Slide credit: Dan Klein and Pieter
Abbeel
<http://ai.berkeley.edu>

Properties of A*

Uniform-Cost



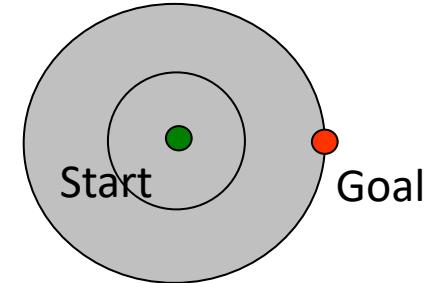
A*



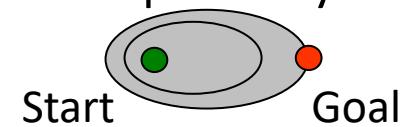
Slide credit: Dan Klein and Pieter
Abbeel
<http://ai.berkeley.edu>

UCS vs A* Contours

Uniform-cost expands equally in all “directions”



A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Slide credit: Dan Klein and Pieter
Abbeel
<http://ai.berkeley.edu>

A* Applications

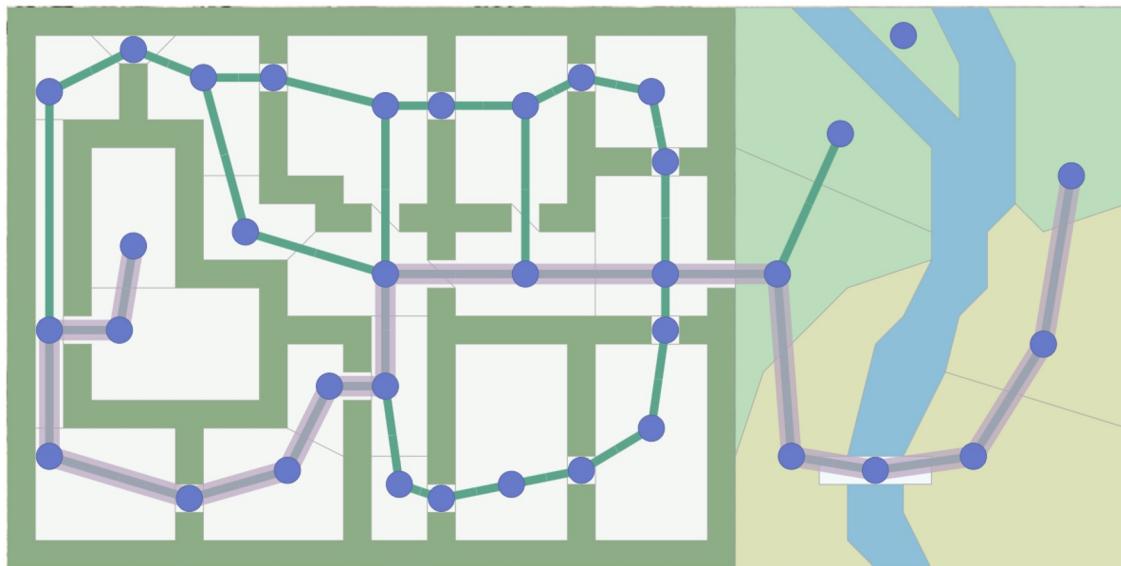
Pathing / routing problems (A* is in your GPS!)

Video games

Robot motion planning

Resource planning problems

...



Supplemental Reading

I recommend this A* tutorial by Amit Patel of Red Blob Games

This is not the solution
to the HW assignment.

<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Introduction to the A* Algorithm

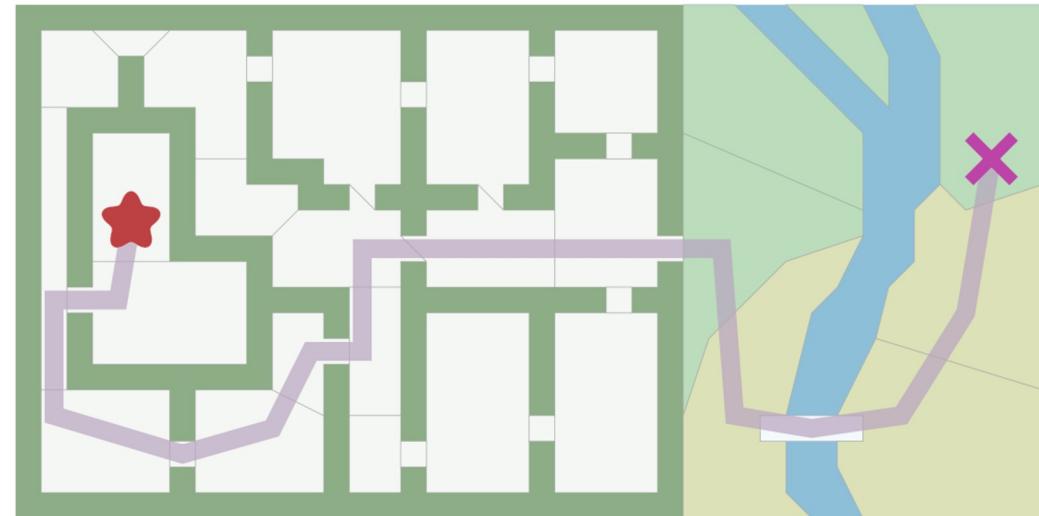
from Red Blob Games

[Home](#) [Blog](#) [Links](#) [Twitter](#) [About](#)

Search

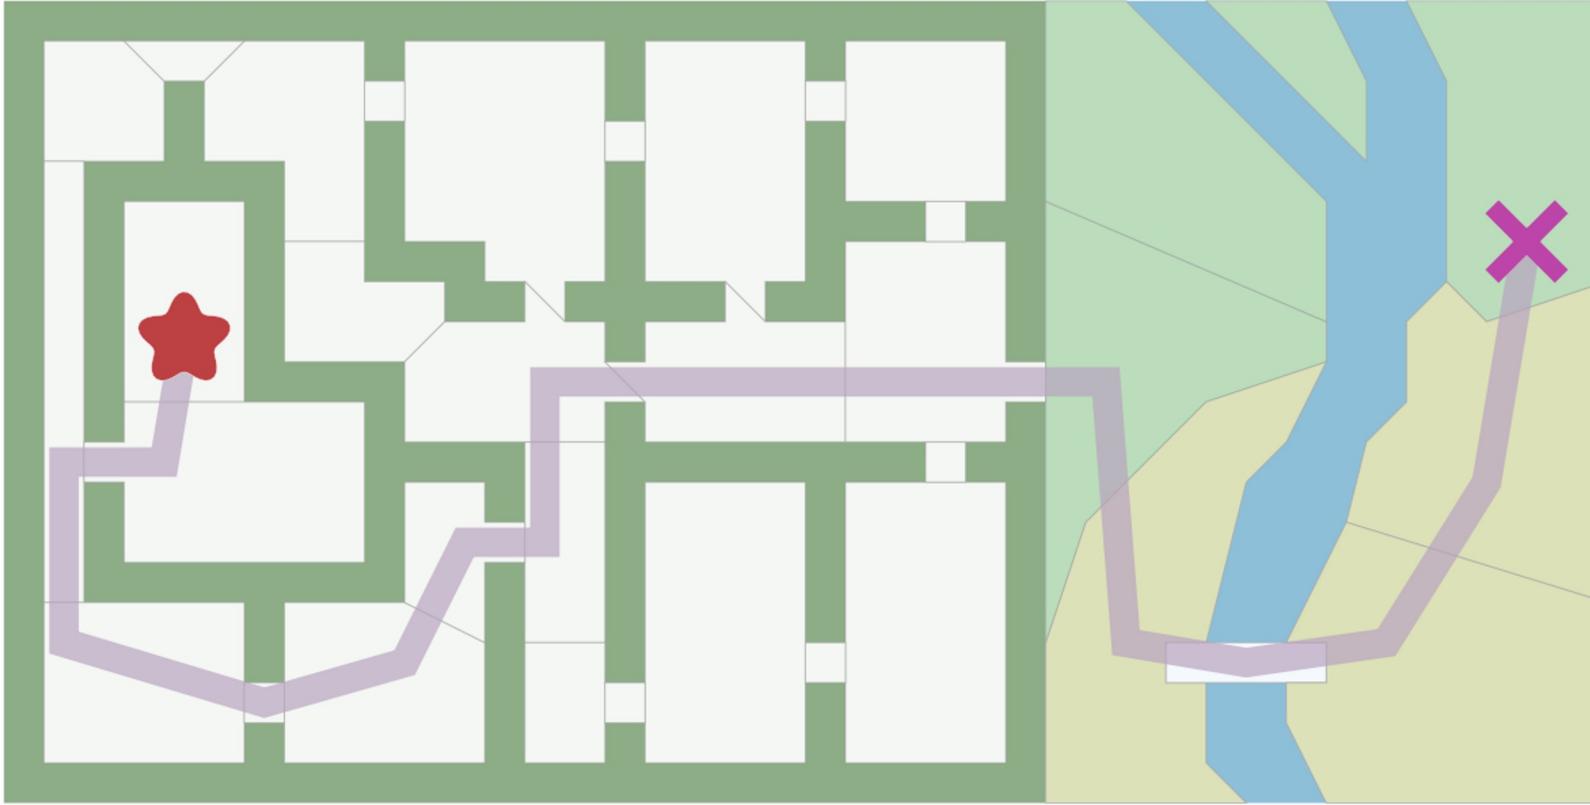
Created 26 May 2014, updated Aug 2014, Feb 2016, Jun 2016

In games we often want to find paths from one location to another. We're not only trying to find the shortest distance; we also want to take into account travel time. Move the blob  (start point) and cross  (end point) to see the shortest path.



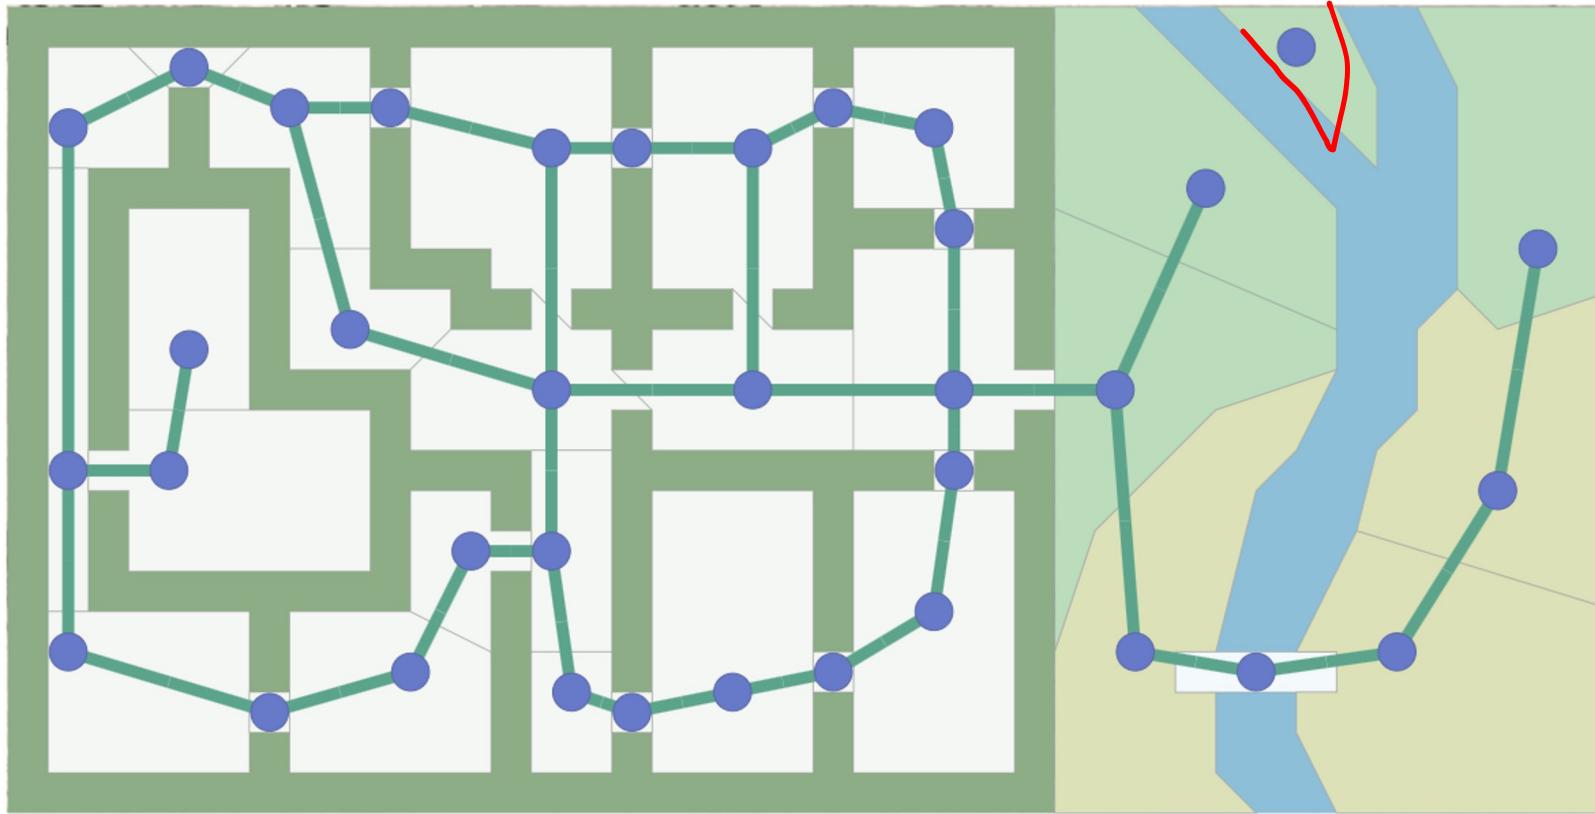
To find this path we can use a *graph search* algorithm, which works when the map is represented as a graph. A* is a popular choice for graph search. **Breadth First Search** is the simplest of the graph search algorithms, so let's start there, and we'll work our way up to A*.

Pathfinding in Games



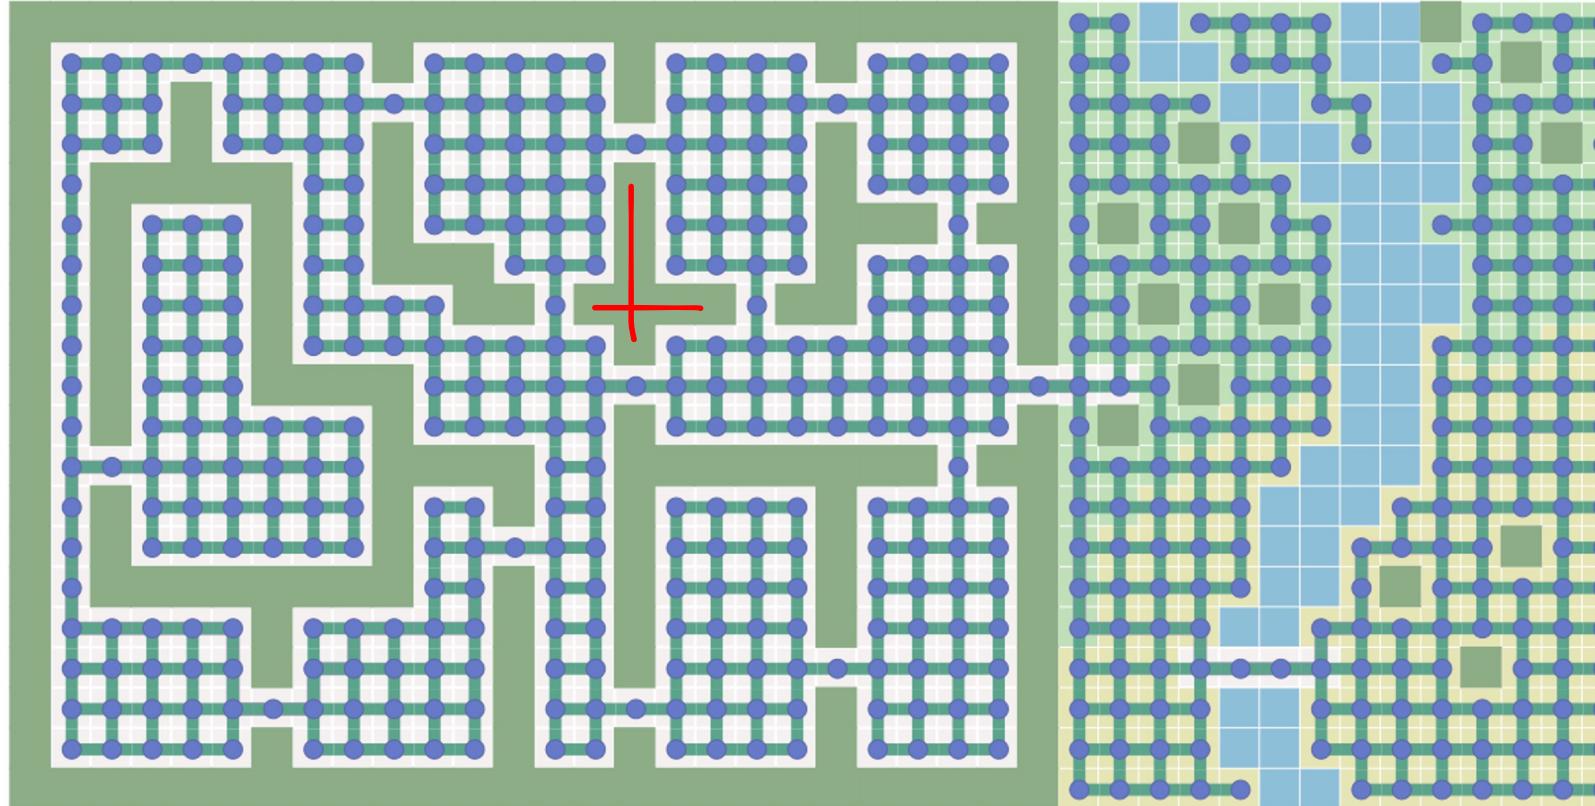
<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Pathfinding in Games



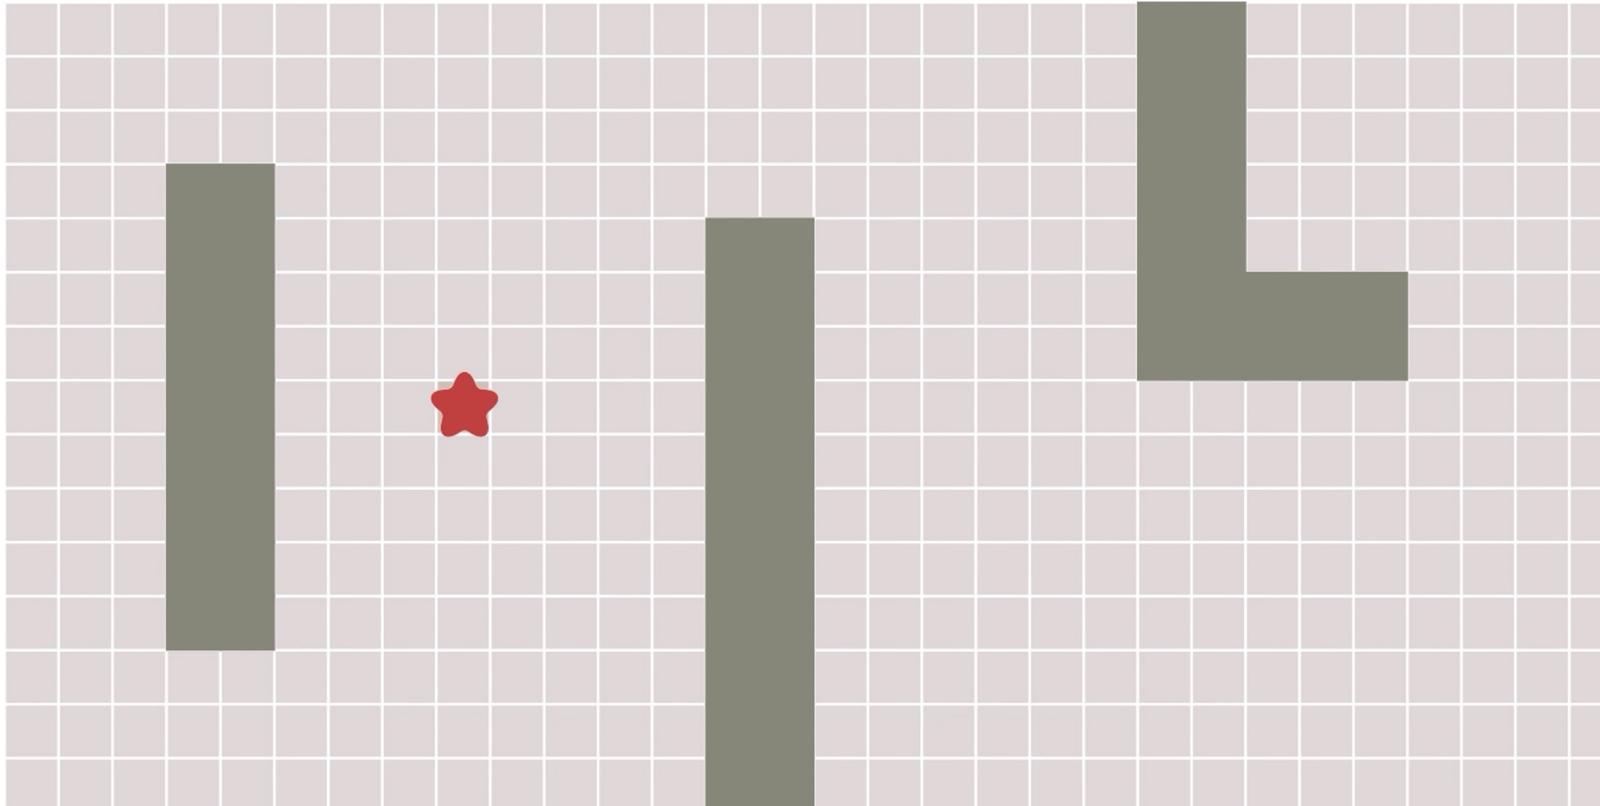
<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Pathfinding in Games



<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Breadth First Search



<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

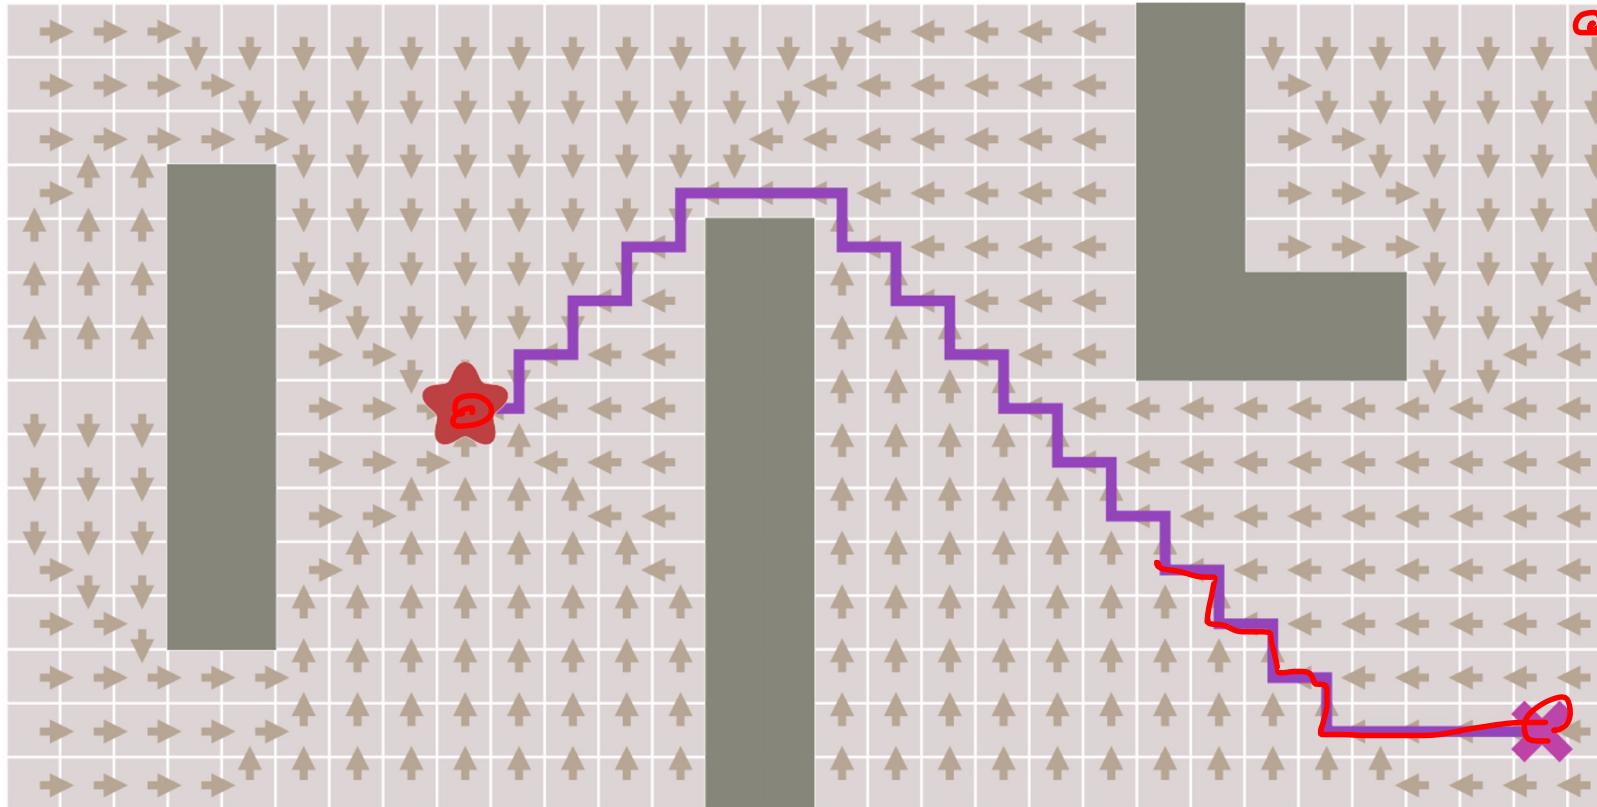
BFS in 10 lines of Python

```
frontier = Queue()
frontier.put(start ★)
visited = {}
visited[start] = True

while not frontier.empty():
    current = frontier.get()
    for next in graph.neighbors(current):
        if next not in visited:
            frontier.put(next)
            visited[next] = True
```

<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Finding the shortest path



<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

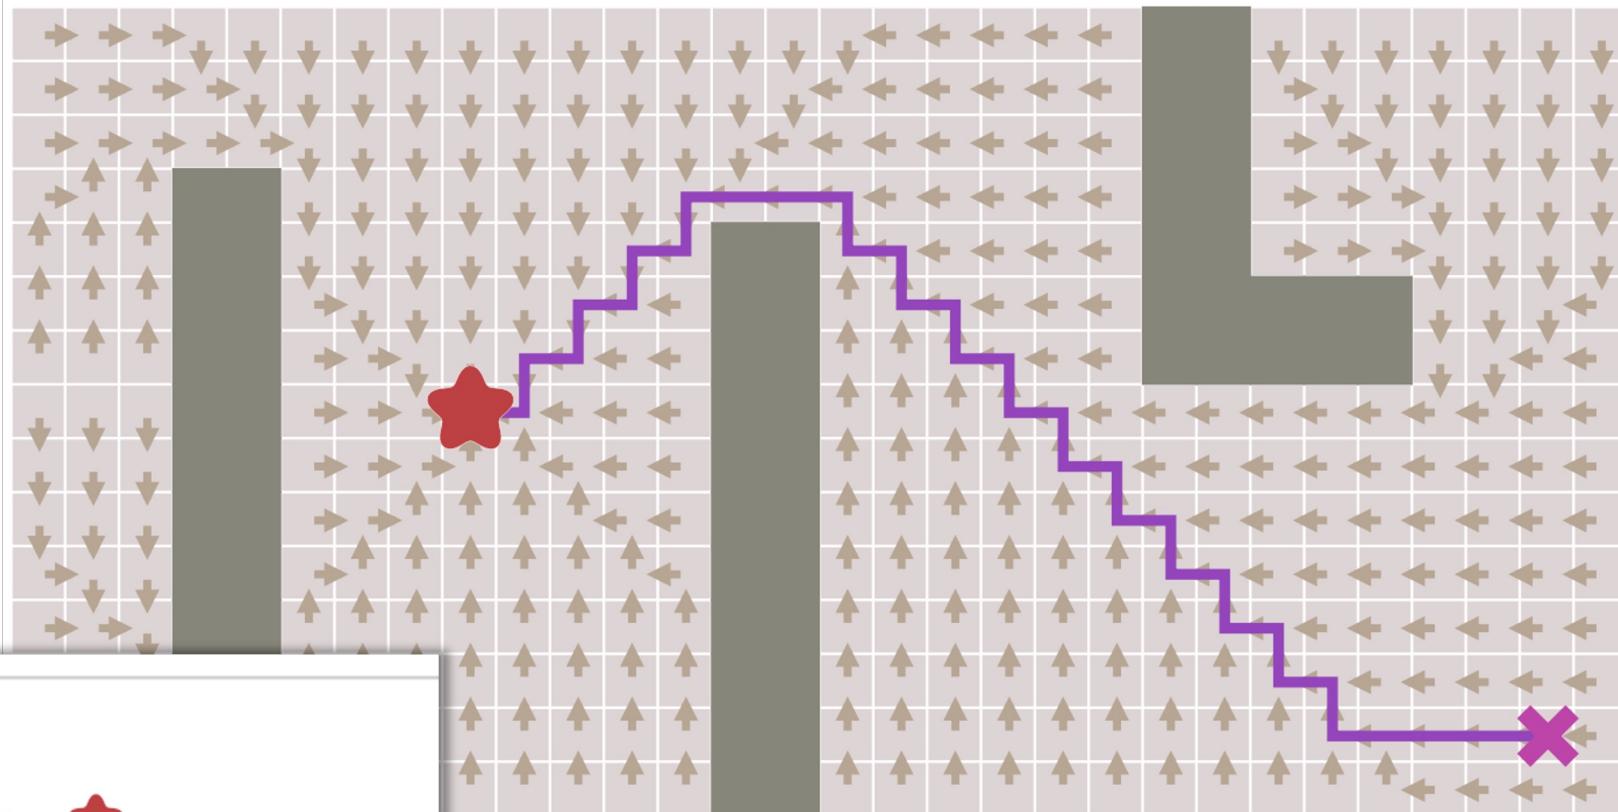
Finding the shortest path

```
frontier = Queue()
frontier.put(start ★)
came_from = {}
came_from[start] = None

while not frontier.empty():
    current = frontier.get()
    for next in graph.neighbors(current):
        if next not in came_from:
            frontier.put(next)
            came_from[next] = current
```

<https://www.redblobgames.com/pathfinding/a-star/introduction.html>

Finding the shortest path



```
current = goal ✎  
path = []  
while current != start: ★  
    path.append(current)  
    current = came_from[current]  
path.append(start) # optional  
path.reverse() # optional
```